



**DIGITAL
RIGHTS
WATCH**

**Submission to the Department of
Infrastructure, Transport, Regional
Development, Communications, Sport,
and the Arts**

regarding the

**Online Safety Amendment (Digital Duty of
Care) Bill 2026**

September 2026

Who we are

Digital Rights Watch is a charity founded in 2016 to promote and defend human rights as realised in the digital age. We stand for privacy, democracy, fairness, and freedom. Digital Rights Watch educates, campaigns, and advocates for a digital environment in which rights are respected, and connection and creativity can flourish. More information about our work is on our website: www.digitalrightswatch.org.au.

Acknowledgement of Country

Digital Rights Watch acknowledges the Traditional Owners of Country throughout Australia and their continuing connection to land and community. We acknowledge the Aboriginal and Torres Strait Islander peoples as the true custodians of this land that was never ceded and pay our respects to their cultures, and to elders past and present.

Statement on the use of AI in this submission

Digital Rights Watch confirms that no generative AI was used in the creation of this submission. Any mistakes herein are genuine human errors and any useful policy ideas are the result of genuine human thought and creativity.

Contact

Tom Sulston | Head of Policy | tom@digitalrightswatch.org.au

Elizabeth O'Shea | Chair | lizzie@digitalrightswatch.org.au

Contents

General Position	4
Proposed Digital Duty of Care	6
Standards for children and adults	6
Harms from generative AI technology	6
Business models of social media platforms	7
User empowerment tools	8
Enforcement	9
Nudification content	10
Transparency	10
Harm definitions	11

General Position

Digital Rights Watch (DRW) welcomes the opportunity to provide feedback on the draft *Online Safety Amendment (Digital Duty of Care) Bill* 2026.

The exposure draft of the bill is an extremely welcome development, and DRW is broadly supportive of this initiative. As per our submission to the 2024 statutory review of the Online Safety Act¹:

“In principle, we agree with the intent of introducing a statutory duty of care for platforms; to flip the burden of responsibility of maintaining safety and security from individuals to corporations who require a social licence to operate in Australia. While we do not see it as a total solution to the problems that the online safety regime is seeking to address, if implemented well it can contribute to improving the situation.

There is an enormous imbalance of power between large technology companies, and the people who use their products and services.”

The Internet is a vital component of modern civil life and it is vital for the health of our democracy and of us as individuals that everyone is able to participate in the online world, free from harms from not only large tech corporations but also regular Australian businesses.

We think it is critically important that a human rights approach to online safety is central and would welcome more explicit commitments in this respect, including in the extraneous materials.

We do have some concerns about gaps in implementation and potential side-effects of some of the proposed legislation, which are addressed in our submission.

¹ <https://digitalrightswatch.org.au/2024/06/25/submission-review-online-safety-act/>

Our main concerns with the exposure draft are:

1. It is not clear how the differential levels of protection for children and adults is justified.
2. Definitions of harm could be updated to better reflect the risks associated with generative AI chatbots and companions, as well as the needs of vulnerable people.
3. The business models of tech companies are not fully contemplated in the framing of the bill, even though they are a significant driver of harm. Social media and other attention-economy business models give rise to the proliferation of extremist content, scams, and deepfake material that is both harmful for individuals and our democracy.
4. The transparency provisions are a good start, but we need greater levels of transparency required of tech companies, including mandating detailed ad libraries.
5. Enforcement options must be improved, specifically to allow the duty to be directly enforceable by people who have been harmed.
6. Removal notices are a start but not enough: there ought to be strict liability for the publication and distribution of nudification content.

Proposed Digital Duty of Care

We welcome the digital duty of care as a positive development for the purposes of creating safer online spaces. Online services make massive profits from our attention and personal information, and there are shockingly few rules about how they do it. We need basic protections to require platforms to remove harmful content and proactively take steps to prevent it from flooding our feeds, including by addressing pecuniary incentives of their own creation, and the proposed digital duty of care advances that cause.

Standards for children and adults

In our view, the differential protection offered to children and adults is not clearly justified. With the exception of pornography, we are of the view that the other material or conduct listed in proposed s 25D also gives rise to harm among adults. We recommend that this be revised, to better protect all users of the online world, and avoid encouraging or requiring the proliferation of unnecessary, unhelpful, and privacy-invasive use of age-gating technology.

Harms from generative AI technology

We propose that the definitions be amended to better cover new and emerging technologies, such as AI chatbots and companion apps. We refer to the Alanah and Madeline Foundation's submission on the form this might take.

We propose that the definitions in this section be amended to capture tools that are not marketed or 'primarily designed' for a harmful purpose, but whose demonstrated use overwhelmingly serves one. Grok's image-generation function is not presented as a tool for producing non-consensual sexually explicit imagery, yet a significant proportion of its documented use involves the generation of child sexual abuse material and non-consensual deepfake sexually explicit material. A definition anchored to stated purpose will be easily evaded by a developer who simply declines to state a harmful purpose. Revealed use, including the composition of a tool's user-base and the content it is actually used to produce, tells a different story.

Specifically, s25A(1)(j)(ii) should be excised as it immediately gives generative AI platforms the ability to opt-out of the duty of care by simply prohibiting (in their terms of service) or removing any social sharing service that they might have, without any requirement to enforce this. The harms caused by the product of their platforms being shared by users will, however, remain.

Business models of social media platforms

It is clear that the underlying business model of social media platforms is a key driver of harm. Social media companies sell advertising space, which is how they generate the vast majority of their revenue. Their competitive advantage is that they monopolise our attention and monetise it, by gathering significant amounts of highly-personal information that allows advertising to be particularly well-targeted.

This results in content appearing on platforms that the platform is paid for (advertising) and content that the platform pays for (such as via influencers). There are very real harms generated by creating a system of pecuniary benefits for attention-seeking content, including that such content is likely to be produced at scale, and be sensationalist, misleading, extremist (or all three) because such content works better with the platforms' approach to microtargeting and amplification.

For these reasons, we recommend that the bill include improved transparency requirements for ad libraries, and endorse the recommendations made in the submission from the Australian Child Rights Task Force. Casting light on these business models, which are otherwise largely concealed from the public, will improve policy discussions.

We also note that the system by which platforms approve ads is typically light-touch; it is clear that there is significant 'inauthentic' behaviour on platforms. Scams are able to be advertised by accounts, including in a coordinated manner, which exploit the minimal oversight exercised by platforms in approving advertising content. Scam content is harmful content, and addressing it should be a requirement to comply with the duty of care.

A related gap concerns the interaction between the duty of care and paid content (such as sponsored posts from individual or business profiles). It is not clear whether advertising posts are captured by the proposed user empowerment measures, including any right to opt out of algorithmic curation or endless scroll. If paid content can continue to be inserted into a feed regardless of a user's feed settings, the practical value of those settings is undermined, since advertisers retain a commercial incentive to override user preferences in pursuit of reach.

We recommend the bill clarify that user empowerment tools apply to *all* content appearing in a user's feed, whether organic or paid, and no matter how it is classified by the platform.

User empowerment tools

We welcome the inclusion of user empowerment tools in the proposed legislative regime, but would welcome greater clarity about how such powers might be used. We think user empowerment tools, including careful use of defaults, are a rights respecting way to give people greater control over their online environments. We would like to see more drafting and clarity around what is proposed, so that this can be properly considered by advocates and specialists in the field.

We note that much of the public discussion of the duty of care has centred around the existence of an “opt-in” or “opt-out” model for algorithmically-generated content feeds. Noting that there is nothing in the duty of care bill that specifically addresses how user consent for design features is gathered by platforms, we must also point out that major social media platforms typically already provide an “opt-out” from an algorithmically-generated feed as they are required to do so under the European Digital Services Act². This has failed to address the issues caused by algorithmic feed generation, as users are disincentivised to use their opt-out through a lack of “sticky” preferences, and there have been cases of malicious compliance. For these reasons, user empowerment tools need to be mandated to default to the most user-protective settings, rather than those that most suit the platform. I.e: opt-in must be the default, with the platforms saving users’ settings so that users are not continually required to assert their wishes.

² <https://digital-strategy.ec.europa.eu/en/policies/digital-services-act>

We note that the bill is silent on empowerment at the community level. Millions of Australians moderate, administer, or otherwise organise online communities, audiences or groups hosted on major social media platforms. These community managers and moderators, whether paid professionals or unpaid volunteers, are typically given little-to-no tooling by platforms to appropriately steward their communities or maintain prosocial norms, and are left to absorb much of the practical burden of managing harmful content that this bill is meant to address.

While we appreciate this is not a bill directly focused on labour protections or the rights of Australian digital workers, these front-line operators currently manage a high amount of harmful content, and, like individual users, would welcome better tools to allow agency over their social governance.

Enforcement

We recommend that a direct right of action be available to individuals who have suffered harm as a result of a failure to comply with the digital duty of care.

Regulator action alone will be insufficient to deal with likely harms. Additionally, the regulator is not in a position to seek or obtain compensation on behalf of affected individuals under the proposed regime. While a complaints regime is important, courts have a real role to play in ensuring the duty is upheld. We note that some of the most significant compensation payments and behavioural changes on the part of social media platforms have come about through litigation in the US that would not currently be available here for a variety of reasons.

One way to begin to rectify this situation is by making it explicit that individuals have the right to sue persons and companies responsible for online services that cause harm as a result of failing to comply with the digital duty of care.

This ought to be a stand alone provision in the act. We note that there are also other options: civil penalties under s 26B could be able to be sought by and awarded to an individual. Alternatively, if a civil penalty proceeding is brought by the Commissioner under s 26D, then an affected individual should be permitted to bring an action to recover compensation for harm

as a result of the conduct that gave rise to the remedial direction. Regimes of this nature already exist in various iterations in statutes like the Fair Work Act, the Privacy Act and the Australian Consumer Law.

Nudification content

Currently, the proposal requires that the Commissioner identify and request the removal of fake nude content. In our view, the penalties for facilitating publication and distribution of fake nude material by online services ought to be significantly stricter. We know from reporting by Indicator Media that mainstream platforms like Meta host significant amounts of advertising content for nudification apps, providing a substantial source of revenue for the company.³ Ads for this technology often target children.

There is no good use-case for nudify apps and fake nudification content. Any online service that hosts it should be strictly liable for civil penalties. If the Commissioner can show that such content was published or distributed, the online service should pay a civil penalty.

Transparency

We welcome the provisions for researcher access to platforms and their data. We defer to academic colleagues for the specifics of how best these arrangements should be structured, noting that the specifics of the bill need to be centered on meeting the needs of researchers, rather than protecting platforms from scrutiny. This would mean that some proposed solutions, such as placing researchers into platform “sandboxes” would be inadequate for quality research and good levels of transparency.

Services that serve advertisements and advertisement-like content such as “promotional posts” need to be compelled to provide greater levels of transparency over the libraries of advertisements that they serve to their users. Currently, platforms do very little moderation of advertisements, and it is not uncommon to see failures of their own policies around political, adult, or harmful advertisements. Given that the whole point of an advertisement is to be seen, we can think of no good reasons why platforms should not be required to hold publicly-accessible and searchable libraries

³ [AI Nudifiers continue to reach millions and make millions](#)

of the advertisements that they have been selling, along with the details of who was paying for those advertisements to be shown.

We are supportive of the additional powers for the eSafety Commissioner under s 26, giving them the ability to demand risk assessments from platforms, offer criticism of those assessments, and directions for remediation.

Harm definitions

We have concerns around the definition of harm, particularly those harms that are not “serious harms”, and the processes by which future harms may be defined. For example, with respect to s25D(1)(d), stories of Ned Kelly could be interpreted as “glorifying crime”, or harm-reduction messaging as part of “harmful practices”. In a context in which harm is defined proscriptively, greater clarity in respect to what does not constitute harm would be welcomed.

Similarly, there are a host of things that are harmful to children that are not named in the legislation, such as the encouragement of gambling, hostile attitudes towards other vulnerable groups in society such as refugees, racial minorities, disabled people, or LGBTI people. While some speech of this nature would likely be a criminal offence, it is not clear that this is the intention of s25C(1)(j) or s25D(1)(d). These provisions are open ended and lack clarity which we are concerned could give rise to applications that do not reflect the intent.

Further there is a clear risk that hate groups in society will attempt to use this legislation to de-platform those that they hate, by claiming that there are potential harms to children. Already, platforms frequently deplatform LGBTI content, or sexual health and education messaging, as they perceive it easier to do so than to contest its validity. Currently, there is a blind spot in the bill, with no counterbalance or right-to-be-published.

Specifically and additionally, we recommend that s25C(2) and s25D(2) be deleted, as it gives the Minister too much power without any rights-based balancing considerations.

Our concerns arise as an outcome of basing the legislation on a “safety” footing, rather than basing the definitions of harm around human rights - both to speak freely and also not to be harmed by others’ speech. Such

considerations could be included in s25H, for example. Without such framing, discussion of harms are grounded in the narrow categories defined in the legislation, rather than in the experience of an individual being harmed. This also means that the legislation will struggle to recognise minor instances of harm to an individual that become serious when aggregated - such as individual messages that might not individually meet a threshold for bullying.

There are similar definitional issues in s25F and s 25G. We believe that these lists of “design features” are intended to be exemplar, rather than exhaustive. However, the drafting is unclear in this matter and should be updated to reflect the existence and inclusion of other harmful design features, including those that have yet to be invented.