



**DIGITAL
RIGHTS
WATCH**

**Submission to the Department of Prime
Minister and Cabinet: Office of AI**

regarding the

National AI Standards paper:

***“Getting it right: Building AI infrastructure that
works for Australia”***

October 2026

Who we are

Digital Rights Watch is a charity founded in 2016 to promote and defend human rights as realised in the digital age. We stand for privacy, democracy, fairness, and freedom. Digital Rights Watch educates, campaigns, and advocates for a digital environment in which rights are respected, and connection and creativity can flourish. More information about our work is on our website: www.digitalrightswatch.org.au.

Acknowledgement of Country

Digital Rights Watch acknowledges the Traditional Owners of Country throughout Australia and their continuing connection to land and community. We acknowledge the Aboriginal and Torres Strait Islander peoples as the true custodians of this land that was never ceded and pay our respects to their cultures, and to elders past and present.

Statement on the use of AI in this submission

Digital Rights Watch confirms that no generative AI was used in the creation of this submission. Any mistakes herein are genuine human errors and any useful policy ideas are the result of genuine human thought and creativity.

Contact

Tom Sulston | Head of Policy | tom@digitalrightswatch.org.au

Elizabeth O'Shea | Chair | lizzie@digitalrightswatch.org.au

Contents

General Position	4
Managing growth in large data centres	5
Planning and community consent	5
Requirements for large data centres	6
Energy and water consumption	6
Resource consumption reporting	7
Conditions for AI training	9
Safety	9
Privacy	9
Copyright	10
Professional responsibility	10

General Position

Digital Rights Watch (DRW) welcomes the opportunity to provide feedback on the *Getting it right: Building AI infrastructure that works for Australia* consultation paper.

DRW recognises the potential of AI to improve Australian standards of living and productivity. However, it also has the potential to cause immense harm to our quality of life, personal safety, and environment.

We welcome the introduction of AI standards into Australia. Not only is this important for creating an environment where future industries can be built, but also for ensuring that those industries are not doing so on the back of infringing Australians' human rights.

Managing growth in large data centres

Planning and community consent

How can communities and industry work together to ensure data centres deliver lasting community benefits? How can upfront commitments be reflected in state and territory planning processes or community engagement?

Data centres need to deliver benefits for local communities, not just the US corporations that profit from them. Currently, the planning and approvals process for AI data centres constitutes a failure of democracy. While federal and state governments are keen to see more data centre construction in their quests for productivity and local infrastructure projects, local councils and communities are ignored in their objections to having their neighbourhoods industrialised with no benefit to local residents.

There needs to be a strengthening of local planning regulations to ensure that residents' voices are heard and centred in planning applications: no data centre construction should proceed without local approval. This needs to be distinct from the centralisation of planning approvals with state administrative tribunals in order to accelerate housing construction. Fast-track approvals and state tribunals over-ruling local planning concerns should be applied only for housing projects, not for data centre construction.

Requirements for large data centres

Energy and water consumption

How should 'new' renewable capacity be defined in relation to data centre investments?

What incentives or obligations would encourage data centres to invest in on-site generation, batteries or other flexible demand measures?

AI data centres require significant and constant consumption of power and water. This is simply the nature of the technology - AI consumes a lot of power because it is extremely computationally inefficient. The servers then require significant cooling and high levels of water consumption as a result.

AI training does not scale linearly with the size of the training data - as the systems become more complex, they will require more power to train. Additionally, as AI inference use increases, the per-user energy demands will also increase. This is the opposite of traditional deterministic digital systems, which experience economies of scale as more users are added.

For this reason, the requirements on data centres need to be strict, well-monitored, and nationally-consistent. We cannot afford the current scenario in which states race to the regulatory bottom in order to attract capital projects.

AI data centres introduce such a huge electricity demand that they must use only renewable energy for Australia to have any chance of meeting our obligations under the Paris climate agreement. That means no use of fossil-fuels for on-site generators and not purchasing any new grid electricity that has not been renewably-generated, even for firming purposes. It also means not simply purchasing carbon offsets.

New renewable capacity needs to be new generation and sufficient storage to allow a data centre to operate under peak demand 24x7, otherwise net fossil-fuel-generated electricity use will increase as a result, with a concomitant increase in carbon emissions.

To allow data centres to consume any fossil-fuels represents a huge risk to Australia - as our economy becomes increasingly dependant on AI, it will also become more dependant on fossil fuels, locking us into high levels of carbon emissions at a time during which we need to be absolutely focused on decreasing them.

Resource consumption reporting

How frequently should data centres be required to transparently and publicly report their energy use, efficiency, and emissions as part of plans demonstrating compliance with mandatory requirements?

The technology exists by which electricity and water usage can be measured and recorded in real time - this already happens for regular Australian households. We see no reason why large data centres couldn't and shouldn't report their usage on a real-time basis. Given their demands involved are huge, this will be necessary data for generation and distribution companies to manage the load.

It will also be necessary for regulators to inspect the use of energy during times of low renewable generation, to ensure that fossil-fuelled firming capacity is limited.

Should a nationally consistent hierarchy of water sources be adopted for data centres, prioritising non-potable/recycled water with potable water as last resort?

The UN has recognised “the right to safe and clean drinking water and sanitation as a human right that is essential for the full enjoyment of life and all human rights”¹, which Australia recognises as part of a right to an adequate standard of living².

¹ <https://digitallibrary.un.org/record/687002?ln=en&v=pdf>

²

<https://www.aq.gov.au/rights-and-protections/human-rights-and-anti-discrimination/human-rights-scrutiny/public-sector-guidance-sheets/right-adequate-standard-living-including-food-water-and-housing>

The question of hierarchy of water sources needs to be reconsidered as a system of rights to those water sources - not only of humans as a biological need, but also as a means by which the natural environment is sustained, and as parts of First Nations country and culture. When considered in this manner, commercial use of water sources is at the bottom of the hierarchy - it is only acceptable for AI data centres to consume water when it does not infringe on our human right to water and a natural ecosystem, as well as our First Nations peoples right to ownership over those water systems³.

³ <https://www.dcceew.gov.au/water/policy/first-nations>

Conditions for AI training

We believe that the adoption of national AI standards offers some important opportunities to set down rules for AI hyperscalers that will keep them from infringing on Australians' human rights for years to come.

In order for any of these rules to be realistically enforceable, they must be grounded in transparency, so that relevant regulators can inspect and assess AI corporations compliance - particularly across the privacy, safety, and copyright domains.

There are many areas across which transparency over an AI corporation are required - from their approaches to testing and safety through to the data sets on which their models are trained.

We also believe that Australia can lead the world in demonstrating that innovation is not constrained by good regulation but enabled by it - consumers of AI products need to know that they are fit for purpose and built in an ethical and sustainable way before they are willing to trust those products.

Safety

Given that AI corporations are currently required to disclose their training data to EU and Californian regulators⁴, this would not produce a significant regulatory burden, but would ensure that it becomes harder for these corporations to use jurisdictional arbitrage to locate models in Australia that use data that is unacceptable in the EU.

For example - the LAION-5B AI training data set contains links to known child sexual abuse material - including images of Australian children being abused. It is unacceptable for AI companies to use this training data; they must be compelled to disclose what training data sets they are using.

Privacy

Given recent developments in privacy legislation, individuals are likely to soon be granted a right to erasure of their personal data. Without

⁴ https://leginfo.legislature.ca.gov/faces/billNavClient.xhtml?bill_id=202320240AB2013
<https://artificialintelligenceact.eu/high-level-summary/>

transparency of data sets used by AI companies, individuals will be unable to even know if they can request deletion from a company using a data set that might contain their data, let alone verify that their request has been actioned.

Copyright

Similarly, copyright holders need to be able to grant or revoke consent for their work to be used in AI models. Transparency over the data used in a model is fundamental to their ability to do so. This has been exploited by AI corporations, which to date have been extremely careless in their attitude to respecting the rights of copyright holders, not just in Australia but globally.

Professional responsibility

AI corporations have recently been responsible for a number of security breaches, where unvalidated AI agents have been exposed to the public internet only to hack into web properties, including those operated by the Australian government.

Given the ability of these agents to rapidly discover and exploit flaws in poorly-secured internet systems, this represents a risk to the safe functioning of Australian businesses and government. While it remains the responsibility of those entities responsible for maintaining websites to keep them secured against both intentional and accidental hacking, there is an important role for AI companies to play when they create agents with hacking capabilities. This is especially relevant when we consider that many of the AI hyperscalers are also US defence contractors. It's not a leap of faith to consider that an AI company training infosec agents in Australia could unwittingly create cybersecurity problems in an unfriendly nation, causing a diplomatic issue for Australia in the process.

Therefore, it is important that companies creating AI agents are bound by not only a legal framework that clearly assigns liability for the actions of those agents, but also a set of professional standards that allow them to demonstrate that they are acting in a responsible manner. These standards should cover issues such as breach notifications, software testing and validation practices, sandbox/airgap definitions, risk assessment and response and mediation plans.

The only way in which compliance with these standards would present a significant burden would be if an AI company were already deficient in their professional standards.

We would consider AISI to be an appropriate entity to define the standards and monitor compliance with them and would like to see the national AI standards issue a directive to AISI to develop and maintain such standards.