AUSTRALIAN COMMUNITY MANAGERS

DIGITAL RIGHTS WATCH

Misogyny in Online Social Spaces

How Community Managers, Moderators and Administrators Govern Gendered Harms

2026

Australian Community Managers and Digital Rights Watch present this report with thanks to the community managers, social media managers and moderators who took part in our survey and attended one-on-one interviews.

This report could not have been produced without generous support from the Minderoo Foundation.

Report design courtesy of gooddept. by Nikita Lamba.

Contents

01 ABOUT THIS REPORT

02 DEFINITIONS

- Communities
- Social media
- Community managers
- Social media managers
- Moderators
- A wider ecosystem

07 RESEARCH DESIGN

- Data analysis

09 THE LANDSCAPE OF MISOGYNY

- Assumed superiority
- Gendered victim narratives
- Manosphere mainstreaming

17 WHEN MISOGYNY SHOWS UP

- Women simply existing
- Women's expertise and authority
- Gender equality and women's rights
- Intersectionality compounds misogyny

23 WHERE MISOGYNY IS HAPPENING

- Open platforms and diminished governance capacity
- Community type and subject matter

27 MODERATION RESPONSES AND PRACTICES

- Phase One: Investigation
- Phase Two: Determining mitigating factors
- Phase Three: Intervention
- Community care
- Moderating toward collective governance
- Record-keeping

Contents

42 EFFECTS OF MISOGYNY

- Social norming
- Chilling effects
- Derailed conversations
- Real-world harms
- Mental health issues for moderators

54 WHAT MAKES IT WORSE

- Under-resourced tools and platform abdication
- Overwhelm
- Contextual complexity
- Lack of understanding and respect
- The online/offline divide
- Structural disempowerment
- Institutional betrayal

65 A SYSTEMIC PROBLEM

68 CONCLUSION

71 NEXT STEPS

76 REFERENCES

78 ABOUT THE AUTHORS

80 APPENDIX A

81 APPENDIX B

83 APPENDIX C

About this report

Misogyny is increasingly recognised as a significant challenge for digital participation, online safety and democratic engagement. While considerable attention has been paid to platform policies, regulation and industrial content moderation systems (Roberts, 2019), the people responsible for its everyday management within the online spaces where it occurs remain far less visible, in both public discourse and decision-making.

Commissioned by Digital Rights Watch and conducted by Australian Community Managers (ACM), this report addresses that gap. Drawing on survey and interview data, it explores how community managers, moderators and social media managers experience, interpret and respond to misogyny across digital social spaces: how it manifests, the strategies used to address it, the emotional and organisational toll of this work, and the conditions shaping governance decisions and practitioner views on what would improve them.

Rather than treating moderation as solely a technical or regulatory challenge, this research approaches it as a social and cultural practice: one that depends on contextual judgement, community knowledge, relationship-building and ongoing stewardship. Practitioners are an essential layer of digital governance: translating platform rules into everyday safety practices, closing gaps where rules are absent, and shaping conditions for constructive, prosocial participation. This is also care work. Practitioners protect others from harm while absorbing much of its cost themselves, and as digital front-line workers, their wellbeing warrants institutional and policy support commensurate with the role they perform.

Two questions anchor this research: what makes online social spaces (both formal communities and more informal transient publics), safe for women and diverse voices to participate in; and what techniques practitioners already use, or could be supported to develop, to build such spaces. The sections that follow address the first by examining how and where misogyny surfaces and the effects it produces. They address the second by documenting the investigative and relational techniques practitioners use to prevent and respond to harm.

By centring practitioner lived experiences and strategies, this report offers stakeholders new insights into how misogyny is encountered and managed everyday in digital social spaces. In doing so, it provides valuable perspectives on the mechanics of community governance and its role in online safety infrastructure.

Definitions

This report refers throughout to communities, social media, community managers, social media managers and moderators. These terms are often used loosely, including within the profession, so it is worth setting out what distinguishes them and how they relate to the wider system of people and processes responsible for governing online spaces.

Communities

Community is a specific social structure, not a platform, technology set, or location.

Online communities are distinct social spaces that form around a shared interest, circumstance or purpose. People in these spaces have common values and a shared sense of belonging. We differentiate online communities from social media on the basis of deeper interpersonal relationships that form in these spaces. This thinking builds on the work of earlier sociologists such as Benedict Anderson (1983) and social psychologists McMillan and Chavis (1986) and their four pillars of 'sense of community": defined membership, influence (or reciprocity) integration and fulfilment of needs and, finally, shared emotional connections. When considering 'sense of virtual community' (Blanchard 2006, 2008, 2011 & 2024) we add the additional pillar of immersion (Koh et al., 2003) and point to the heightened need for trust in online space (Blanchard and Markus, 2004). Figure 1. below gives a visual interpretation of this understanding.

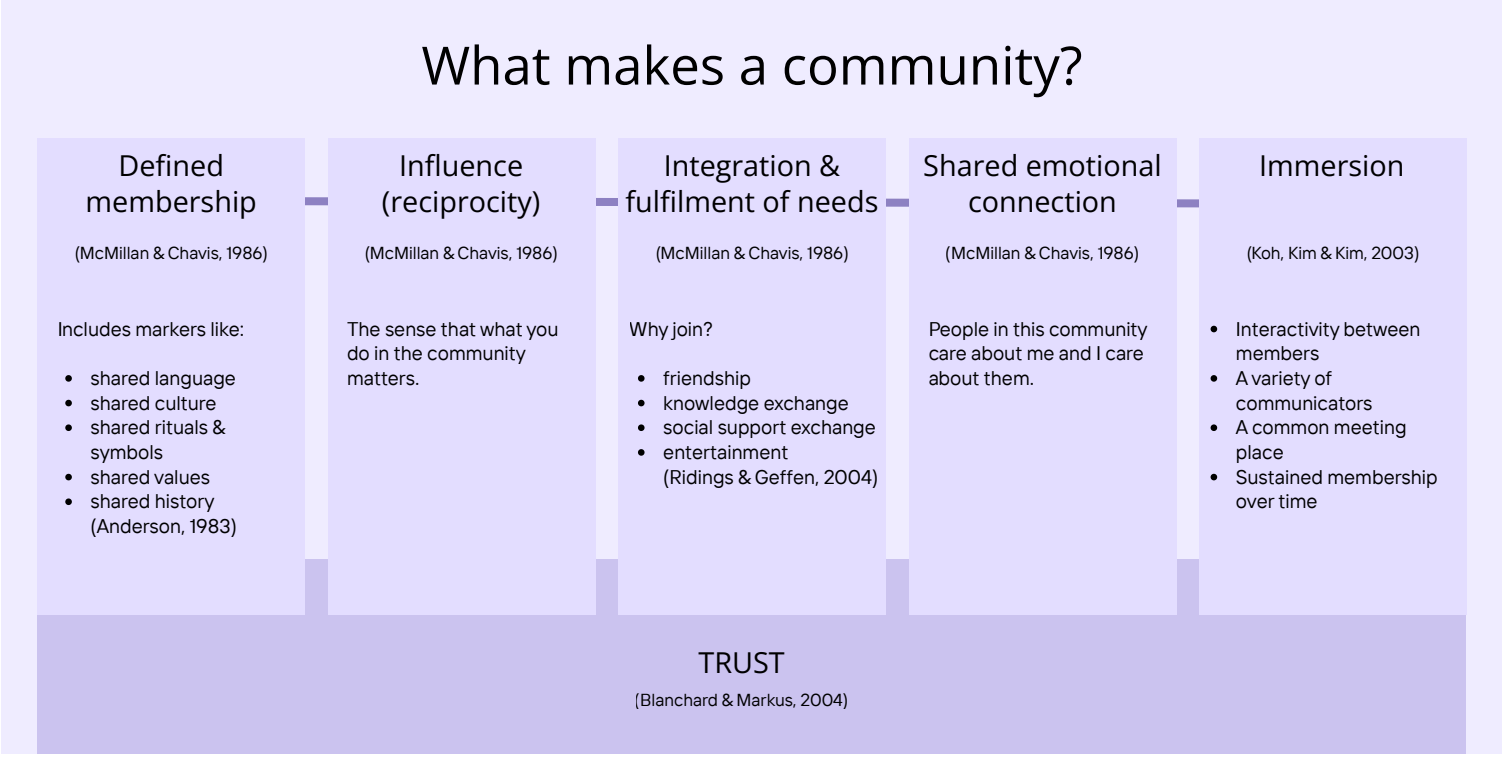

Figure 1: Defining elements of community

Social media

Social media refers to a set of digital platforms and technologies that enable users to create, share and interact with content, and to communicate with others online.

Social media accounts are generally run for the benefit of a business or organisation, or serve as people's personal spaces. Structurally, they are closer to broadcast channels that allow and often encourage commentary, rather than spaces built for sustained exchange between members.

This distinction can be understood through Granovetter's (1973) theory of weak ties. Social media connections are typically weak ties: low-commitment, low-intimacy relations that circulate information efficiently but carry little of the reciprocity or emotional investment McMillan and Chavis (1986) identify as central to community. A shared interest in a brand may create the appearance of community, but it rarely produces defined membership, mutual influence or shared emotional connection between participants themselves. Interactivity remains largely one-to-many rather than between members (Jones, 1997), and the trust, immersion and belonging that mark a virtual community (Blanchard, 2006, 2008, 2011, 2024; Koh et al., 2003) are correspondingly diminished.

In practice, social media accounts are used for promotion, brand building and customer service. This distinction matters to the findings in this report: it is precisely the weak-tie, broadcast character of social media, versus the strong-tie interactivity of community, that shapes the tenor and type of moderation practice.

Community managers

Community managers build and steward online or hybrid groups over time.

The role centres on relationship-building, norm-setting, welcoming new members and facilitating participation. Moderation sits within this labour: it is a routine act of care and governance, woven into the broader work of cultivating a healthy group culture.

Social media managers

Social media managers run branded or organisational accounts and channels.

They focus on content, audience growth and public-facing engagement.

Moderation is usually part of the job: responding to comments, actioning reported content and enforcing house rules. It can compete with content and marketing priorities.

Moderators

Moderators apply rules and norms to content and behaviour:

reviewing, actioning and intervening for norm and boundary management. This can be a dedicated role in its own right, paid or volunteer, sitting apart from community or social media management. It can equally be folded into either of those roles as one part of a larger job.

Across all three, moderation is best understood as a discipline and a practice; rather than a job title or series of tasks. It is a behavioural intervention, applying a norm to a specific piece of behaviour or content, that produces a cultural outcome: the tone, safety and health of a social space over time. Online community managers moderate as standard practice, social media managers usually do, and a dedicated moderator role can also stand alone. The discipline is constant even where the job title changes.

None of these roles is tied to a single employment model. Community managers, social media managers and moderators can be paid or volunteer, and the work can sit in-house, with an agency, or with an independent freelancer or consultant. This variation shapes the resourcing, authority and job security practitioners describe throughout this report.

A wider ecosystem

Community managers, social media managers and moderators sit within a larger moderation ecosystem.

Commercial content moderators, typically employed by or contracted to major platforms, review flagged or reported content at scale, usually without any relationship to the community whose content they assess. This labour is infamous for its precarity, traumatising impacts and proximity to modern slavery (Pinchevski, 2023; Perrigo, 2022; Newton, 2019; Roberts, 2019).

Platform trust and safety teams design and enforce the policies and systems that govern entire platforms, working at least one level removed from any single social space. Regulators, such as Australia's eSafety Commissioner, set the outer legal boundary within which all of this activity sits. Community managers, social media managers and moderators occupy the layer closest to the community itself, applying a depth of context and relationship the layers around them cannot.

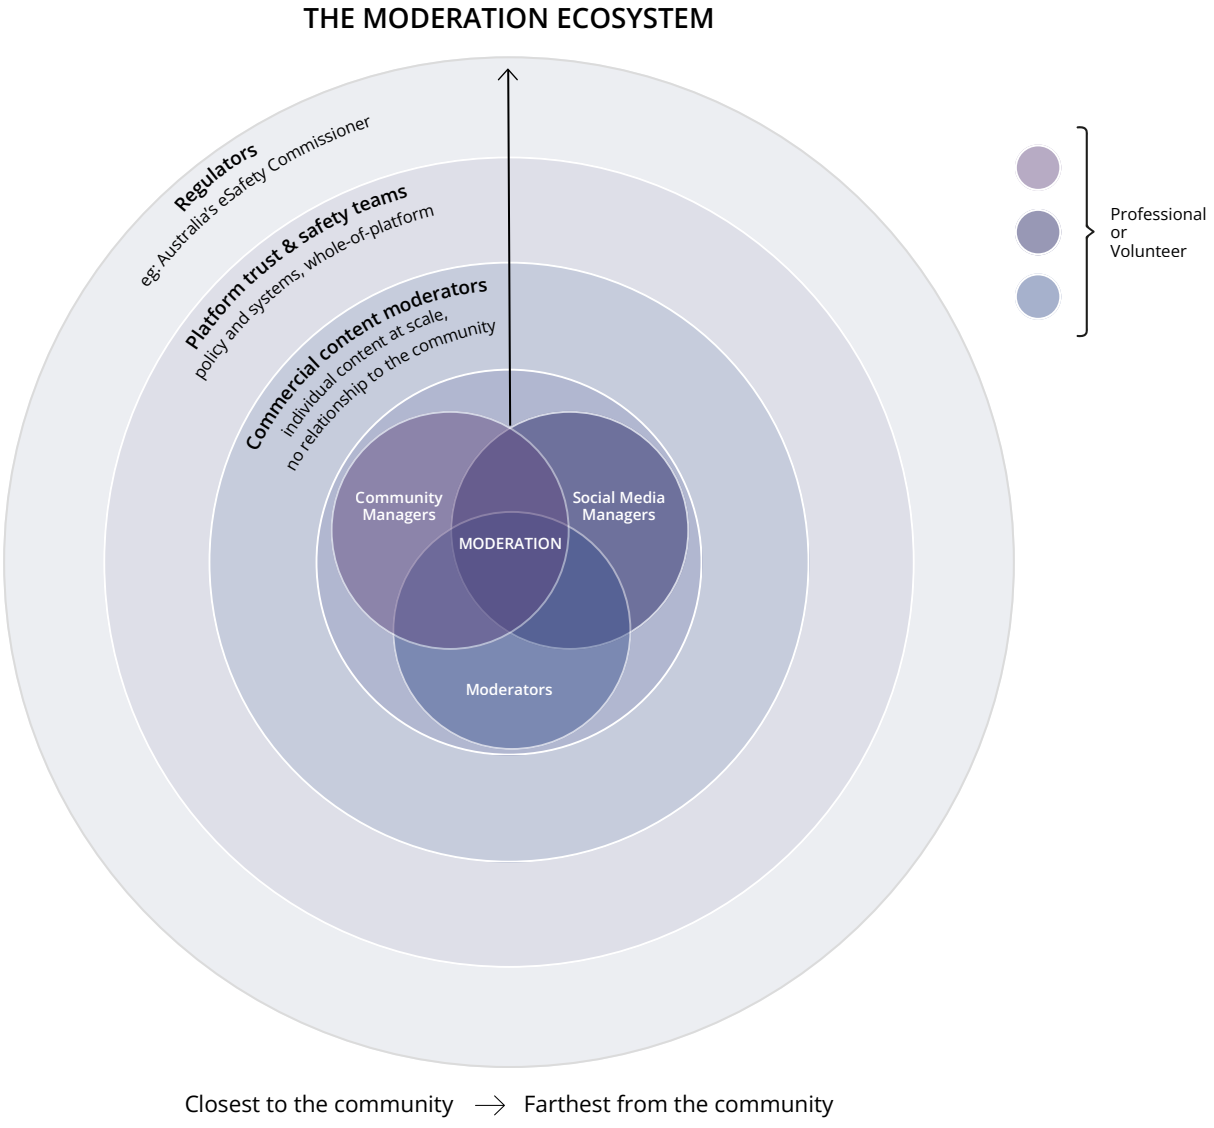

Figure 2. The moderation ecosystem, from regulators to the community itself, and the spectrum of employment across community-facing roles.

Research design

Both authors write from an insider standpoint. Dr. Venessa Paech has worked in online community management and moderation for three decades, and Dr. Jennifer Beckett has also worked professionally in the field. Both currently train and support community managers, social media managers and moderators. This proximity shaped the research from question design through to interpretation, providing access and fluency unlikely to be available to an outside researcher. This perspective influenced the design of the research and interpretation of results. Following feminist standpoint epistemology (Hartsock, 1983), we treat this as a situated perspective to be made explicit rather than a bias to be neutralised.

This report draws on two primary sources of data: a survey of Australian community managers, moderators and social media professionals, and a series of semi-structured interviews with practitioners. Participants came from a wide range of sectors including mental health, financial services, retail, non-profit and community services.

The survey ran from 1 April to 28 June 2026 and was promoted through ACM channels, direct outreach by the authors and snowball sampling, and received 43 responses. Respondents were predominantly professional practitioners whose work involved some level of moderation (58% professional community managers, 19% professional social media managers and 5% professional moderators), with a smaller proportion working in volunteer roles. Most were highly experienced, with 70% having worked in the field for eight years or more and managed a wide range of community spaces and/or social media accounts. These spaces were most commonly: corporate or brand (56%), advocacy, activist or feminist (30%), local government or community services (28%) and news or media organisations (14%).

Interview participants were recruited from survey respondents who volunteered for a follow-up interview. Twelve people were selected based on their experience and to ensure an even spread across types of online social spaces and industries. Of those 12, six agreed to be interviewed. Prior to the interviews, all participants received detailed and plain language information statements, provided informed consent, and were reminded of their right to pause, skip questions or withdraw at any time. Participants were also offered the option of a pseudonym, which all six elected to use. Employer names and identifying operational details were removed or generalised in interview transcripts and throughout this report to protect participant confidentiality. Anonymised participant summaries are provided in Appendix C.

A standard interview protocol was developed between the two researchers in order to ensure commonality across interviews. Interviews covered a range of topics from how and when misogyny turns up in the online spaces participants managed, to its effects and how it is managed. Interviews were conducted over Zoom as participants came from across Australia and the globe, providing more flexibility for participants. Interviews were recorded using the in-built audio recording feature on Zoom and an initial transcript was created using Zoom's transcription feature. Transcripts were cleaned by the researchers to ensure accuracy and were then sent back to participants for review. Participants were given two weeks to respond with any revisions before analysis took place.

Data Analysis

Transcripts and qualitative survey responses were uploaded into nVivo for analysis.

> Inductive coding was used as it allows for codes and themes to emerge from the data allowing meaning-making as it is compared with current theory and literature (Bingham, 2023). Researchers coded the data together through a process of constant comparison and internal memoing, allowing for consistency in analysis. During the initial open coding phase of the analysis codes were placed under broad categories as they emerged, these were later audited and some codes were combined and/or moved into different themes.

The landscape of misogyny online

Practitioners in this research consistently described misogyny as a routine, rather than exceptional, feature of digital sociality. Across professional associations, gaming communities, advocacy organisations, local government, news organisations, brand communities and health or help-seeking spaces, respondents encountered misogynistic behaviour as part of the ordinary work of digital community governance and moderation. They situated it within a wider social and political context rather than treating it as an isolated phenomenon.

Participants characterised misogyny as cumulative, adaptive and context-dependent. They described the spaces they govern as sites where prevailing cultural narratives are reproduced and amplified. Effects included discouraged participation, undermined credibility, derailed conversations and a narrowing of who feels able or willing to contribute. Many linked an apparent rise in misogynistic behaviour to the current political climate, which they felt had granted greater licence to express views once left unspoken; several also noted that online discourse now supplies ready-made talking points and rhetorical strategies. Practitioners thus framed misogyny as behaviour shaped and legitimised by a wider discursive environment.

Our coded analysis identified three interrelated themes that practitioners consistently returned to when describing this broader environment: assumptions of male superiority, narratives of male victimhood, and the mainstreaming of ideas associated with the online 'manosphere' (Sugiura, 2021). Together, these themes provide an important lens for understanding both individual acts of misogyny and the broader cultural conditions within which community managers, social media managers and moderators are expected to govern.

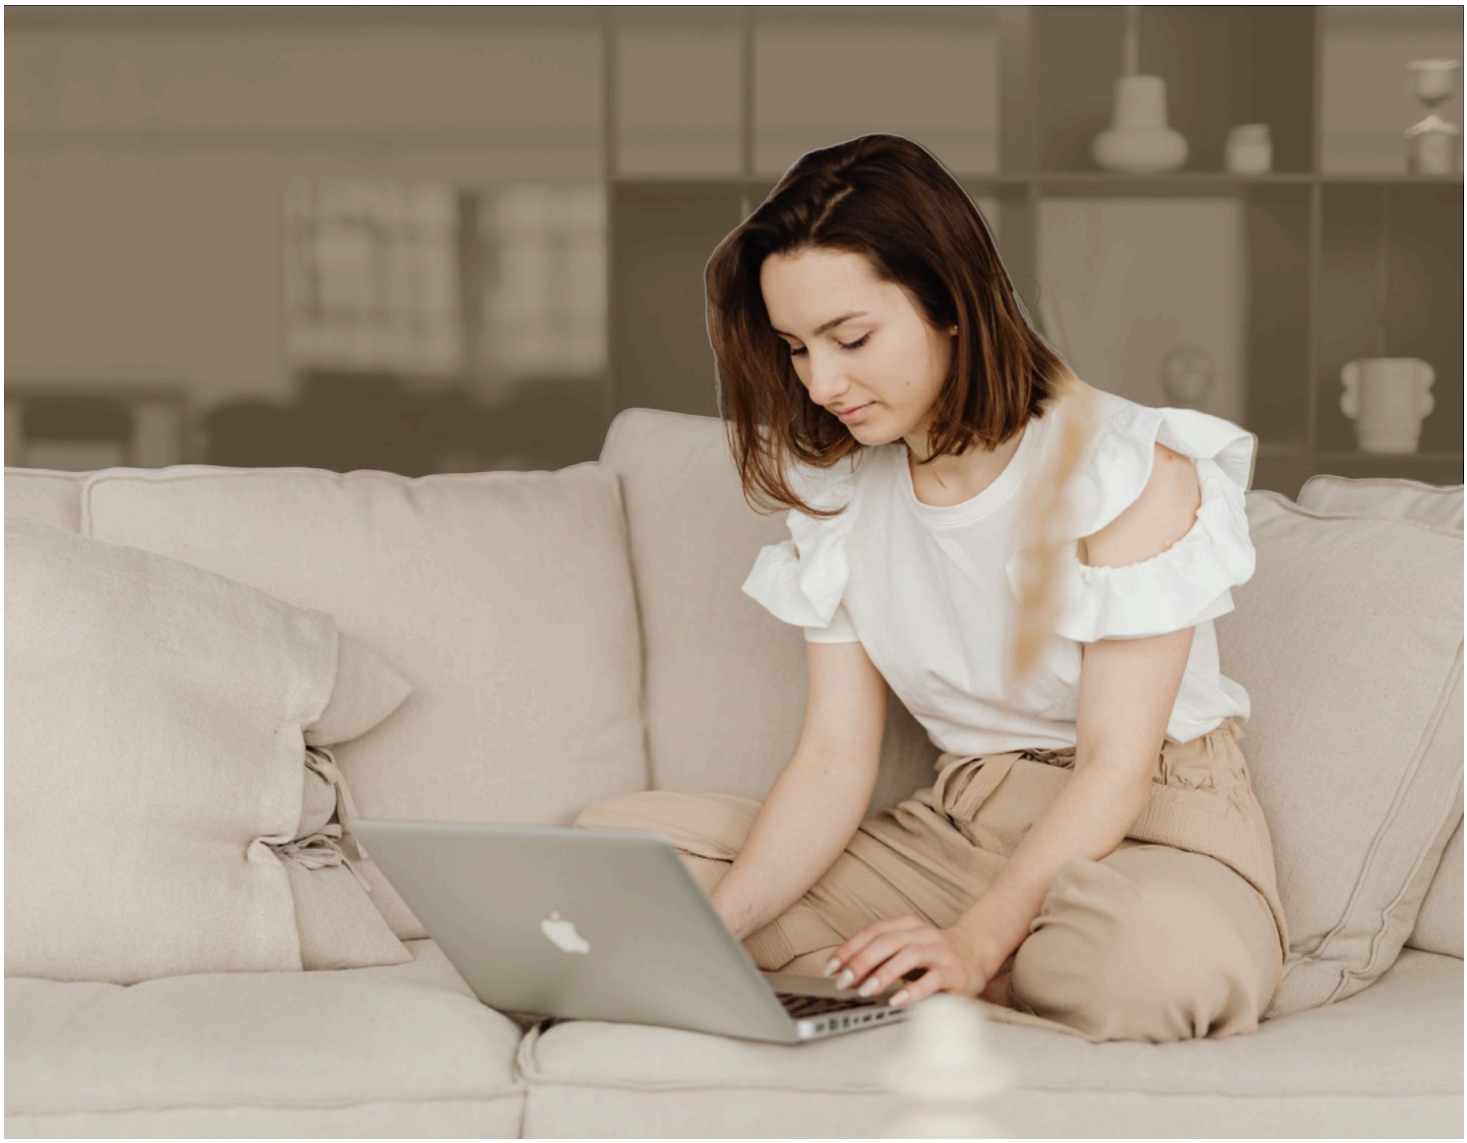

Assumed superiority

The most pervasive theme identified by participants was the assumption that women occupy a less legitimate position within online spaces.

Respondents rarely described misogyny as random hostility. Instead, they observed behaviours that sought to reinforce traditional gender hierarchies by questioning women's expertise, authority, leadership or right to participate.

This pattern emerged consistently across the diverse online spaces represented in the study, although it manifested differently depending on whether practitioners were governing audiences, publics, groups or communities. Although these behaviours took different forms, respondents understood them as expressions of the same underlying dynamic: challenging women's legitimacy to occupy public and communal space.

Importantly, practitioners did not identify these patterns solely through individual comments or pieces of content. Rather, they described recognising recurring forms of behaviour that unfolded across conversations, interactions and relationships. It was the cumulative pattern—who was targeted, whose authority was questioned, when hostility escalated, and how other participants responded—that signalled misogyny. This distinction emerged repeatedly throughout this research and is one of this report's central findings: practitioners are rarely moderating individual posts in isolation. Instead, they are interpreting behaviour within its broader social context, drawing on historical knowledge, community memory and an understanding of how exclusionary dynamics develop over time.

One practitioner with more than two decades' experience across public institutions, higher education and the private sector described this assumption as close to universal:

> “There’s a bias to imagining the person that you’re interacting with is male. If you don’t know their gender, you’ll assume male and straight. I feel that communities are largely set up with men in mind.
>
> - FRANKIE
> *social media and community management practitioner*

Gendered victimhood narratives

Alongside assumptions of male superiority, practitioners consistently identified narratives that repositioned men as the primary victims of efforts to address gender inequality.

These narratives frequently emerged in response to discussions about women's leadership, domestic and family violence, workplace inclusion, women's health, gender equality initiatives and women's safety.

Respondents reported that conversations centred on women's experiences were regularly interrupted or reframed through familiar rhetorical devices. Comments such as "what about men?", claims that men now face greater discrimination than women, or assertions that initiatives supporting women unfairly disadvantage men were described across audiences, publics, groups and communities alike. Participants rarely interpreted these interventions as genuine attempts to broaden discussion. Rather, they described them as recurring discursive strategies that redirected attention away from women's experiences while reframing conversations around perceived threats to male status or privilege. As one agency-based community manager, who moderates across a range of client sectors, put it after fielding a men's rights comment on an unrelated post: “no, patriarchy sucks for everybody, actually, you too, turns out.”

For practitioners, these exchanges presented a particularly complex governance challenge because they rarely involved obvious policy breaches. Viewed in isolation, many comments appeared civil, relevant or framed as legitimate disagreement. Their significance emerged only when considered within the broader dynamics of the space in which they occurred: how they redirected conversations, discouraged participation, accumulated over time and reinforced exclusionary norms. Respondents therefore described relying heavily on contextual judgement to distinguish robust disagreement from behaviour that functioned to silence, derail or marginalise women's participation. Once again, participants emphasised that effective governance depended less on evaluating individual comments than on recognising the behavioural patterns and social effects those comments produced collectively.

Manosphere mainstreaming

Perhaps the most striking finding emerging from the qualitative analysis was practitioners' perception that ideas traditionally associated with the online manosphere were increasingly appearing within mainstream digital spaces.

Participants described encountering narratives, language and behaviours that echoed those found within men's rights communities, incel forums and other anti-feminist online networks.

Our participants noted that rather than relying on overtly misogynistic language, users increasingly drew upon shared talking points, memes, rhetorical frames and coded expressions from the manosphere that could be adapted to a wide variety of contexts. They frequently observed that these narratives appeared remarkably similar across otherwise unrelated online spaces, suggesting that misogynistic discourse was no longer confined to particular subcultures but had become part of a broader digital vernacular.

One participant, who manages social channels for a major superannuation fund, described watching this mainstreaming happen across a single membership base over several years:

> “Anytime any of our marketing used imagery of women, someone would comment and say, ‘why is she a plumber, not making my sandwich?’ We used to laugh at it because it’s such old humour. Now it’s less of that and way more sophisticated—you see the Trump supporters, the Tate supporters. Definitely more varied and aggressive too.
>
> - PRIYA
> *social media manager, superannuation sector*

Several interview participants situated this shift within a wider political and cultural context. They described an environment in which increasingly polarised public debate, influential media figures, online commentators and algorithmically amplified content had created greater social licence to express misogynistic views publicly. Significantly, this environment provided individuals with an expanding repertoire of ready-made narratives, rhetorical strategies and conversational scripts that could be readily deployed across different online settings. Participants therefore described contemporary misogyny as operating not only through individual prejudice but through the circulation of shared discourses that legitimise, normalise and reproduce gendered hostility across digital life.

For practitioners, this mainstreaming fundamentally altered the nature of governance. Whether managing a transient social media audience, facilitating a public discussion, stewarding a membership community or moderating an ongoing group, respondents increasingly found themselves interpreting irony, humour, coded language, dog whistles and patterns of interaction whose meaning depended almost entirely on context. The challenge was therefore not simply recognising offensive content, but recognising when seemingly ordinary interactions formed part of broader patterns of exclusion, intimidation or harassment. This distinction between moderating individual pieces of content and governing patterns of behaviour runs throughout the findings that follow and helps explain why respondents consistently viewed human judgement, contextual interpretation and accumulated community knowledge as indispensable to effective moderation.

When misogyny shows up

Participants described misogyny as pervasive across the online spaces they governed, and non-random. Across both the survey and interviews, practitioners consistently identified particular situations, topics and forms of participation that reliably attracted misogynistic responses. Over time, many described developing an ability to anticipate where misogyny was likely to emerge, allowing them to prepare moderation strategies, allocate resources and intervene before discussions deteriorated.

Survey data supports this picture of pervasiveness. Asked how often they encountered misogynistic content or behaviour in the spaces they managed, 44% of respondents said regularly or constantly, and a further 42% said occasionally; only 14% said rarely. Asked how misogyny typically showed up, the most common form by a wide margin was generalised sexist commentary or "jokes" (79%), followed by targeted abuse directed at specific individuals (47%) and subtle or coded content that was hard to classify (44%). One in five respondents (21%) reported that misogynistic content had been directed at them personally, in their capacity as moderator, admin or manager, rather than at other community members.

This anticipatory knowledge was particularly evident among experienced practitioners. Rather than viewing incidents in isolation, respondents described recognising recurring patterns that cut across audiences, publics, groups and communities. While the specific topics varied depending on the nature of the online space, the underlying dynamics remained remarkably consistent. Misogyny was most likely to emerge when existing gender norms were perceived to be challenged, or when women simply occupied visible space online.

The data identified clear recurring triggers.

Women simply existing

The clearest finding across the dataset was that the presence of women alone was often sufficient to provoke hostility.

Several survey respondents captured this finding with striking simplicity:

"The presence of women is enough."
"Anything that features a woman honestly."
"Any content that features a woman..."

Rather than identifying controversial topics, respondents repeatedly identified the **presence of women** as the common denominator.

Across the dataset, practitioners described misogyny emerging around:

- women appearing in photographs or videos
- women featured in promotional campaigns
- women's programs
- women's achievements
- women in leadership
- women in male-dominated professions
- women speaking publicly
- women expressing expertise
- women participating in online discussion.

This pattern emerged across news organisations, advocacy organisations, local government, professional associations, technology communities, gaming spaces and commercial social media channels. Whether practitioners were governing large audiences, publics, membership groups or ongoing communities, they described women's presence itself as a recurring catalyst for misogynistic engagement.

Importantly, the expression of misogyny varied according to the social formation being governed.

Public-facing audiences on social media (usually characterised by more ambient interaction) were more likely to produce:

- comments about women's appearance
- sexualised remarks
- objectification
- criticism of women's public visibility

Groups and communities, where participation and relationships are ongoing, more commonly produced:

- challenges to women's expertise
- questioning of women's authority
- undermining of women's legitimacy
- attempts to exclude women from participation

Although behaviours differed in form, respondents interpreted them as reflecting the same underlying dynamic: attempts to regulate women's presence and participation in digital life.

For experienced practitioners, this pattern had become sufficiently predictable that it shaped everyday governance practice. Rather than simply reacting to misogyny after it occurred, practitioners described developing an expertise in recognising the conditions under which it was likely to emerge. In particular, they noted anticipating increased moderation activity whenever women featured prominently in content or discussions, allowing them to prepare moderators, monitor conversations more closely and intervene earlier if necessary.

Women's expertise and authority

Beyond women's visibility, practitioners identified women's authority as a significant trigger for misogynistic behaviour. Respondents reported that hostility intensified when women occupied positions associated with leadership, expertise or decision-making; particularly when that expertise was in spaces traditionally associated with men.

Common examples included:

- women in leadership roles
- women sharing professional expertise
- women working in STEM
- women in trades
- women in policing and emergency services
- women in sport
- women discussing barriers they had experienced
- women moderating or governing online spaces themselves

Participants frequently observed that equivalent contributions made by men rarely attracted the same degree of scrutiny. Instead, women were more likely to have their qualifications questioned, their motives challenged or their authority undermined.

For practitioners, these behaviours were significant because they were often subtle rather than overt. Individually, many comments appeared relatively innocuous. Collectively, however, they produced environments in which women's expertise was persistently contested, discouraging participation and reinforcing unequal power relations within digital spaces.

Gender equality and women's rights

Participants also consistently identified discussions relating to gender equality as predictable flashpoints for misogynistic engagement.

Topics repeatedly mentioned included:

- feminism
- domestic and family violence
- gender equality
- the gender pay gap
- sexism
- women's safety
- reproductive rights
- relationships and dating dynamics

Rather than engaging with the issues themselves, practitioners described conversations being redirected through familiar rhetorical patterns.

These included:

- "What about men?"
- accusations of bias
- accusations of being "woke"
- denial of structural inequality
- interventions from men's rights perspectives

Respondents rarely interpreted these responses as isolated disagreements. Instead, they described them as recurring conversational strategies that diverted attention away from women's experiences while reshaping discussions around perceived threats to male status.

For moderators, these exchanges presented particular governance challenges because they often sat within the boundaries of acceptable debate when viewed individually. Their cumulative effect, however, was to derail conversations, discourage participation and shift community norms over time.

Intersectionality compounds misogyny

While misogyny was experienced across a broad range of online spaces, practitioners consistently observed that it became more intense when gender intersected with other marginalised identities. This created layered experiences of abuse that were often more severe for users and moderators.

Participants most commonly identified heightened misogynistic responses in discussions involving:

- trans women
- women of colour
- First Nations people
- LGBTQIA+ communities
- feminist advocacy
- broader social justice and inclusion initiatives

Respondents noted that hostility directed towards these groups was rarely confined to a single form of prejudice. Instead, attacks often combined multiple identities simultaneously, with gender becoming one element of broader campaigns of exclusion or harassment. Practitioners described moderation practices and processes (explored in more detail later in this report) that required them to interpret overlapping forms of harm rather than treating misogyny, racism or transphobia as separate governance issues.

Advocacy organisations and issue-based publics were particularly likely to encounter coordinated hostility directed at feminist, LGBTQIA+ and anti-racism initiatives, while broader social media audiences more commonly experienced intersectional abuse within comment sections responding to individual women or public figures. Across these different contexts, practitioners consistently emphasised that understanding the interaction between multiple forms of discrimination was essential to recognising the nature and severity of the harm being experienced.

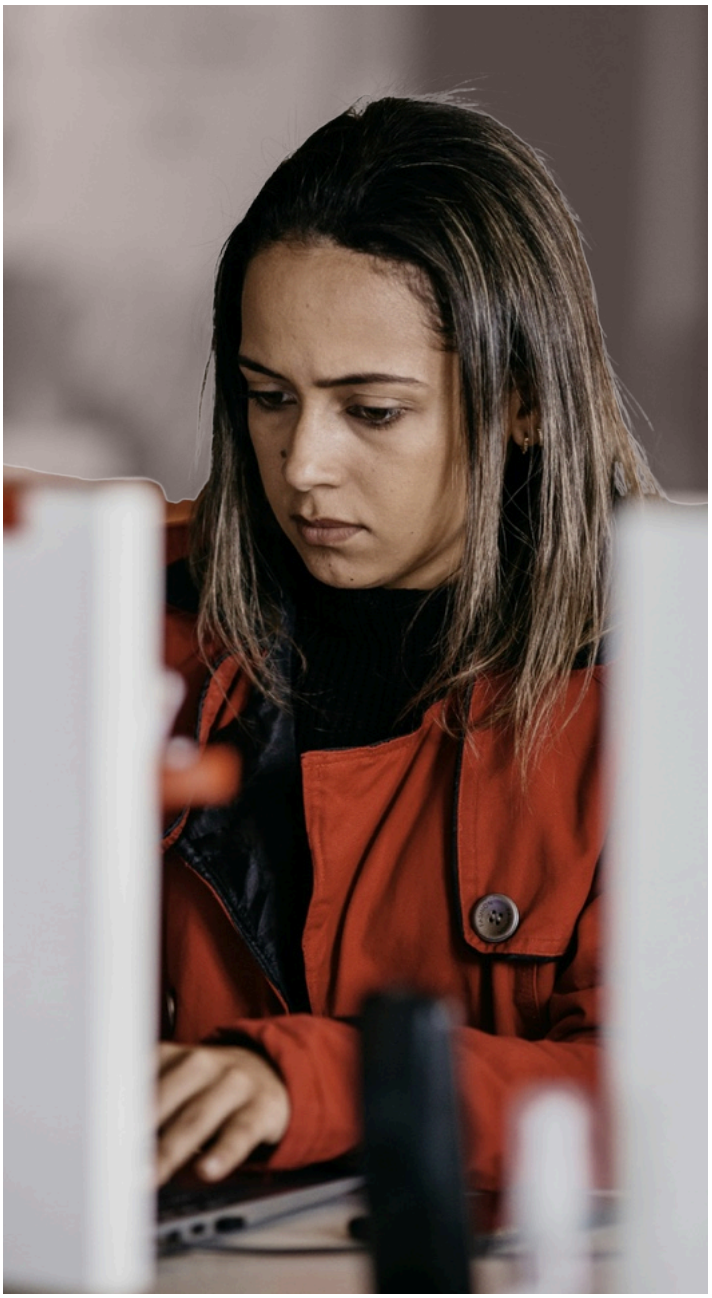

Importantly, respondents suggested that moderation frameworks—and, in some cases, platform policies—often struggle to account for these intersecting harms. Individual comments may appear to target gender, race or sexuality in isolation, yet practitioners understand and must govern the cumulative experience as something far more complex, requiring deep contextual governance.

Where misogyny is happening

Where misogyny surfaced, and how often, was shaped by platform architecture, governance model and the character of the community itself. Practitioners rarely described misogyny as evenly distributed across their work. Instead, its intensity was highly contextual to the specific community or brand being governed, and to the degree of control practitioners held over that space.

Open platforms and diminished governance capacity

Respondents managing open, low-moderation platforms such as X, described more frequent and aggressive misogynistic behaviour than those working within closed or owned communities.

On these platforms, norms are set largely outside practitioners' control, and enforcement depends on external reporting systems over which localised moderators have little influence.

One practitioner at a large organisation described the futility of trying to moderate X:

> ...you can't moderate that channel because you can't delete other people's posts, like you can on Facebook or Instagram. You have to rely on their reporting function, and if you report something to X, they don't take hate commentary off anyway. It's just a pointless exercise.
>
> - SURVEY RESPONDENT

Closed and owned communities were not free of misogyny, but practitioners there were able to set membership criteria and behavioural norms directly (from onboarding users to ongoing boundary maintenance), could deal with moderation issues directly at their discretion, and could intervene or remove participants without waiting on a third party platform's response.

Several survey respondents also linked the character of the misogyny they experienced to platform demographics. Older-skewing users on Facebook were said to engage in some of the most entrenched, least self-correcting misogynistic behaviour, while 'younger' platforms (such as Discord) served up more violent and aggressive forms of it.

Community type and subject matter

The character of misogyny also varied with the type of space being governed.

Gaming and technology communities showed the strongest reactions when women were perceived as experts or challenged established norms, particularly in discussions of sexism within the sector or diversity initiatives.

News and media environments proved especially vulnerable to pile-ons and comment-section misogyny, with strong reactions to coverage of domestic violence, women's achievements, political stories involving women, and celebrity women or women in public life.

> Anytime we talk about the gender pay gap in super, there are always comments that pop up: 'it's a myth, women have it better than men already.' International Women's Day is something we over-prepare for every year, because we know we're going to get that commentary.
>
> - PRIYA
> *social media manager, superannuation sector*

Local government and community services most often encountered backlash around women's programs, community engagement initiatives for women, family violence discussions and inclusion programs, frequently framed by respondents as variations of "why isn't there an equivalent men's program?"

What stayed constant across these contexts was the trigger, not the form: misogyny surfaced whenever women's visibility, authority, or advocacy for gender equality were foregrounded, or sometimes even simply raised, within the specific audience and purpose of the space in question.

Moderation responses and practices

Once a possible instance of misogyny had been identified, practitioners moved through a broadly consistent, sequential process before deciding how, or whether, to intervene: first investigating what a piece of behaviour actually was, then weighing a range of mitigating factors, and only then selecting an intervention. Respondents described each phase as shaping the next; acting on assumption rather than investigation, or intervening before weighing context, was consistently described as leading to worse outcomes.

Phase One: Investigation

This report identifies this first phase as a practice of *contextual forensics*: respondents rarely evaluated isolated comments or posts in isolation.

Instead, they drew on historical knowledge of particular users, familiarity with the online culture of their social space, and awareness of misogynistic narratives and tactics in the broader landscape, to distinguish disagreement from gender-based hostility, and isolated incidents from coordinated harassment.

A number of respondents described actively monitoring spaces associated with manosphere movements to track evolving language, narratives and coded expressions; a form of investigative work that platforms themselves rarely undertake.

> ...that kind of comment triggers my spidey senses, and it might make me go look at their profile a bit further, see if there's been any other similar commentary, or try and ascertain if something elsewhere on social media is driving that sentiment. It's a bit of a forensic deep dive.
>
> - TOM
> *community and corporate affairs practitioner, retail sector*

Where commerical content moderators engaged by platforms, or platform automated tools, typically evaluate individual pieces of content in isolation, practitioners evaluate behaviour within a localised social context. They interpret risk and meaning, identifying and acting on harms that remain invisible to third party operators or automated moderation systems.

74% of survey respondents described themselves as very confident in identifying misogynistic content, and a further 26% as somewhat confident. Not a single respondent selected “not very confident” or “unsure”.

Techniques employed:

- Deep dives into user profiles, posting history and cross-platform activity
- Independent research into language, symbols and coded terms circulating in manosphere-adjacent spaces
- Assessing whether a single instance forms part of a broader, cumulative pattern rather than an isolated event

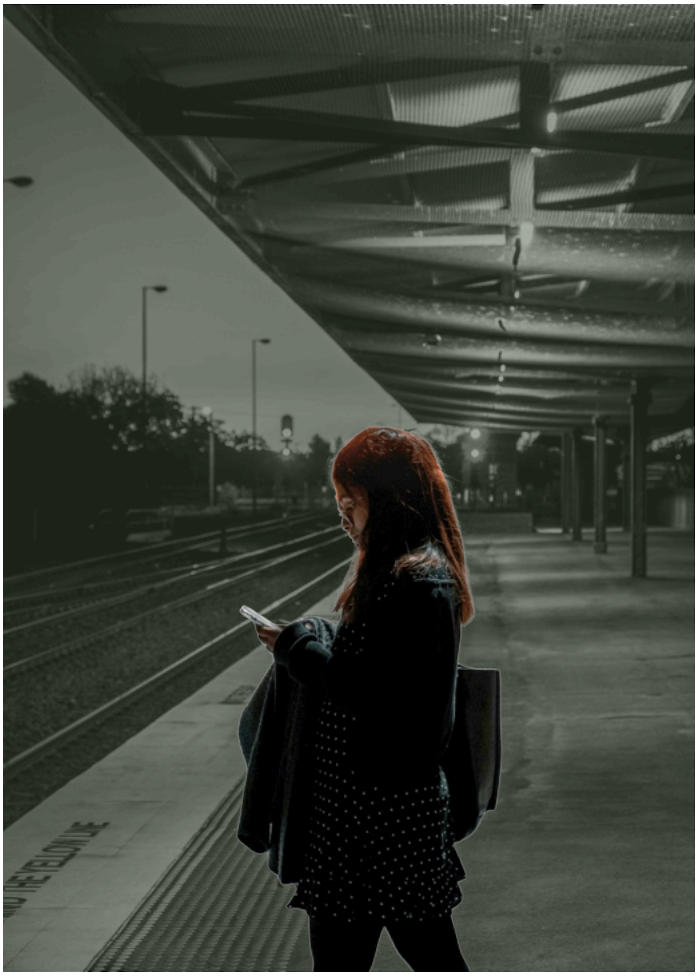

Phase Two: Determining mitigating factors

Once behaviour was identified as potentially misogynistic, the response that followed was shaped by factors beyond the incident itself.

The availability and quality of institutional support proved decisive. Respondents working within not-for-profit organisations, volunteer-led communities and advocacy movements frequently described operating with limited resources and underdeveloped governance frameworks, relying heavily on personal judgement where formal policies or resources were absent or vague.

More significantly, respondents described tensions between community safety and organisational risk (especially in corporate settings): when incidents escalated to senior leadership, decisions were sometimes reframed around reputation, legal liability or stakeholder relationships, creating pressure to minimise, delay or avoid action, particularly where the person involved held influence, donor status or organisational standing.

This pattern closely matches what psychologist Jennifer Freyd has termed institutional betrayal: wrongdoing perpetrated by an institution against the people who depend on it, including its failure to prevent harm or to respond supportively once harm has been disclosed (Freyd, 2008; Smith & Freyd, 2014).

Individual capacity and vantage point also shaped response. Self-identified male respondents in this sample showed a high degree of awareness of misogyny and its impacts, but tended to frame it primarily as a governance and behaviour-management challenge, emphasising rule enforcement and moderation decisions. Self-identified female respondents more often described its emotional and social consequences, including for themselves as practitioners.

The following were key factors survey and interview respondents described as informing their technique and approach:

- Institutional support levels
- Organisational appetite for risk
- Organisational policies and remit of role (e.g. agency practitioner working for a client with strict guidelines, versus one who instructs them to rarely intervene)
- Individual practitioner capacity and bandwidth (workload, pressures, mental state)
- Presence and effectiveness of community guidelines (e.g. Can they be relied on as a touchstone in this instance?; Do they have gaps?)
- Socio-cultural norms active in the social space
- Social contract with users (Is it established? Consitently upheld?)
- Standpoint from which the behaviour is being read: i.e. whose safety and comfort weighed against whose freedom of expression?

Survey responses confirmed that few practitioners rely on any single source of authority to make this decision. Asked how they usually decided how to respond to misogynistic content, the largest group by far (72%) reported using a combination of policies and personal judgement, well ahead of those who relied primarily on formal policy (42%), their own judgement alone (42%) or community guidelines exclusively (56%). Note that response categories were not mutually exclusive.

Only 16% said they typically consulted a colleague or supervisor, suggesting that even when moderators in a team or peers are accessible, the weighing of mitigating factors happens individually, in the moment, rather than through formal processes.

One moderator working in a mental health peer-support setting described this variation vividly, holding up a small gap between her fingers to illustrate it on camera during her interview:

> Our windows of tolerance... mine's like [a small space]. Obviously it's my job, it has to be wider than that, but it makes a difference in terms of how we might respond to something.
>
> - PAIGE
> *online engagement team leader, mental health support forum*

Phase Three: Intervention

At the level of individual incidents, respondents drew on a familiar set of tools:

- keyword filters and automated word blocking
- removing content and behaviour that breaches guidelines
- hiding content from general view without full removal
- turning off comments on specific posts or threads to curb damage
- account blocking and banning
- reporting systems for clear policy breaches

A sub-section of respondents described a preference for hiding content over deleting it, primarily because it avoids escalating the interaction:

> It's just no one's harmed with a little hide, you know? The person still thinks they got their message across without knowing that no one else could see it.
>
> - PAIGE
> *online engagement team leader, mental health support forum*

A preference for hiding over deletion or escalation is best understood as a form of self-care and community care as much as a moderation choice: it helps avoid the pushback, the "why did you silence me" escalation, and the direct confrontation that deletion or removal can provoke, protecting both the practitioner and the wider community from a drawn-out conflict.

However it is a limited remedy. In Australia at least, hiding a comment does not reduce an organisation's legal exposure in the way removal can, since the comment remains published and, depending on platform and settings, still visible to the person who posted it, to their networks, or to anyone with a direct link. The comment itself continues to circulate on the platform rather than being removed from it. Hiding is best understood as a micro-quarantine rather than a resolution: it protects the immediate community from seeing further harm, without addressing the harm's presence, or its potential legal and reputational exposure, at the source.

These were generally viewed as effective for overt misogyny and explicit slurs, but as a first line of defence rather than a complete response.

Prevention over punishment

The more significant finding was practitioners' emphasis on prevention over punishment.

> Rather than waiting for behaviour to escalate, respondents described redirecting conversations early, reinforcing expectations, privately warning users and increasing moderator visibility before norms deteriorated.

Most described favouring education and restorative approaches over immediate exclusion, including one-to-one conversations, moderating "out-loud" within the space to make care visible and enable broader social learning, explanations of community expectations and graduated pathways that reserved exclusion for repeated or deliberate violations (e.g. a known bad actor).

This approach, more consistently available within owned or closed communities than on open, algorithmically mediated social platforms, has recently been packaged by the trust and safety field under the banner of prosocial moderation (Grüning & Kamin, 2025), presented as an emerging alternative to punitive, reactive enforcement. But community management has practised prosociality as default for decades: stewardship, facilitation and restorative conflict resolution have been part of the discipline's working vocabulary since long before "prosocial," a term Wispé (1972) coined in social psychology, entered platform governance discourse. This alignment is nonetheless welcome: if trust and safety and community management take the time to learn from each other, decades of practitioner experience could sharpen a model the field is only now formalising.

Participants in this study described a consistent set of techniques underpinning this proactive and pre-emptive approach:

> - naming the specific impact/s of a comment or behaviour, rather than simply citing the rule it breached;
> - private, one-to-one conversations that invite reflection and a path back to good standing, rather than immediate public sanction;
> - graduated response ladders, reserving removal or exclusion for repeated or deliberate violations;
> - activating bystanders and allies, modelling the response practitioners want to see so other members feel able to speak up themselves;
> - warm handoffs to relevant support services, particularly in mental health and peer-support contexts, rather than removal alone;
> - pre-bunking where misinformation around gender was involved.

These techniques are most effective where there is a sustained relationship with users over time, which favours owned or closed communities with a stable membership over open social platforms, where a harmful actor may never return and bystanders are harder to identify or activate. Importantly, respondents were clear that education and dialogue apply to ambiguous, first-time or lower-severity behaviour, while threats, coordinated harassment and repeated violations still warrant direct removal or escalation, with formal reporting where appropriate.

Several respondents were also careful to distinguish restorative practice from asking the target of misogyny to extend understanding or absorb further emotional work on behalf of the person who harmed them. In their accounts, it is the practitioner, not the target, who carries the burden of engagement.

One agency-based community manager, who moderates for a client whose mission includes directly challenging misogynistic attitudes, described the effectiveness of active, prosocial pushback, and its rarity in her professional settings:

> If we start to push back, you will get other people coming in and supporting, who now feel like they can say something back. We've given them permission. But for a lot of my clients that kind of engagement isn't desired. I wish more people employing community managers would be comfortable with that approach, because all the evidence shows it's way more effective.
>
> - ROBYN
> *agency-based community manager*

These practices are, of course, more labour-intensive: engaging directly with someone who has caused harm, rather than simply removing their comment, is itself a form of emotional labour that requires institutional backing, time and psychological safety practitioners do not always have. For professionals, commercial pressures and imperatives often directly collide with these best practices.

Community care

Underlying these choices is a form of community care that is often invisible in itself: practitioners described weighing not only how to respond to a piece of misogynistic content, but how its visibility, or its absence, would be experienced by the people they represent.

Leaving hateful material up, even briefly, means members absorb it directly; acting on it visibly and promptly signals that the space is being actively stewarded and that members' safety is being taken seriously. Many participants treated this visible governance as itself a protective intervention and spoke of how this approach had led to less instances over time as perpetrators learned the space was not OK with the behaviour.

This created a tension for practitioners who understood the value of visible, active governance but did not always control the standards applied to it. Agency-based and in-house practitioners alike described being directed by clients or organisations towards a "lighter touch" approach, minimising visible moderation for reasons ranging from protecting brand image to avoiding perceived overreach, even where practitioners' own training, research and lived experience indicated this was less effective at deterring harm and reassuring affected members. Robyn's account above, of clients uncomfortable with visible pushback despite the evidence for its effectiveness, was not unusual. Where practitioners lacked the authority to apply the level of governance they believed a situation warranted, the result was often a widening gap between what they knew to work and what they were permitted to do, a further expression of the structural disempowerment discussed elsewhere in this report.

Moderating toward collective governance

A majority of respondents across both the survey and the interviews understood and expressed their work as holistic governance, rather than a more 'simplistic' notion of removing undesirable content.

They discussed the broader work of shaping social conditions under which misogyny becomes less likely to emerge and less likely to be tolerated. Practices include, but are not limited to: articulating behavioural expectations, modelling respectful participation and reinforcing inclusion, seeding and reinforcing constructive social norms, and supporting self-regulation, so users feel empowered and safe to challenge misogynistic comments, support targeted individuals and reinforce community expectations without practitioner intervention.

One veteran practitioner described the latter as a slow, deliberate act of cultivation:

> Building a healthy community from the start is the biggest thing you can do to get self-moderation off the ground... It's like growing a tree—you plant the seed, you nurture it really carefully. Over time, those communities can grow in a way that self-corrects really well.
>
> - FRANKIE
> *social media and community management practitioner*

This governance work extended to the direct protection of targeted members, not only the management and sanctioning of perpetrators. Respondents described checking in with affected members, offering reassurance and support, connecting them to further resources and reducing their exposure to harms, practices largely absent from platform moderation frameworks, which focus on isolated content removal.

Record-keeping

Record-keeping surfaced repeatedly across the survey and interviews as a core technique in managing misogyny.

Contextual forensics depends on record-keeping directly: recognising a coordinated pattern instead of an isolated comment, or a repeat actor instead of a first-time one, relies on practitioners being able to retrieve what happened before (particularly if working with team members across channels or periods of time).

Respondents described keeping their own informal logs of prior incidents, warnings and escalations (e.g. personal spreadsheets, private notes and screenshots) often because no other record existed to draw on and formal tools were inadequate or unavailable. There are exceptions: one interviewee described how their third party tool was a "life-saver" as a record keeping and reporting tool around moderation.

This documentation also carries a regulatory dimension: community managers and social media managers are often expected to produce evidence of what occurred, when, and what action was taken, whether for an internal escalation, a platform report or, occasionally, a legal or regulatory request, placing a compliance-adjacent responsibility on practitioners that is rarely acknowledged or resourced as such. Australia's Online Safety Act (2021) also requires that organisations hosting interactions on social media (not just platforms themselves) be able to produce moderation data if compelled (such as the nature and volume of issues, and details of responses), in a reasonable timeframe.

The same records that support day-to-day moderation decisions are also the closest thing this field has to a real-time evidence base of how gendered harms actually manifest and evolve online, information of direct value to researchers, platforms and policymakers seeking to understand these dynamics rather than infer them. When that record exists only in one person's inbox or memory, it is lost the moment they leave a role.

Effects of misogyny

The consequences of misogyny in these online spaces extend beyond the individual comments or incidents that prompt moderation. Respondents described effects operating at three levels: on audiences and communities themselves, on the practitioners who govern them, and on wider society.

Social norming

Practitioners consistently described misogyny as self-reinforcing.

Left unaddressed, individual instances of hostility signalled to other participants that such behaviour was acceptable, encouraging further instances and gradually normalising conduct that would otherwise breach community expectations. This social norming effect meant that early, visible intervention carried significance beyond the incident itself: it demonstrated to the wider audience or community what would and would not be tolerated.

One agency-based community manager described how quickly this dynamic takes hold once a single comment is left unaddressed:

> If you allow it through, it can turn into a pile-on. People see that someone said a really crook thing and think, time for my hot take. They see one and think they can say whatever, and that applies to everything, racism too. What people see is what they imagine is acceptable.
>
> – ROBYN
> *agency-based community manager*

Chilling effects

Where intervention did not occur, respondents described a chilling effect on participation: women and other targeted members withdrew, posted less, or left the space altogether rather than continue to absorb hostility.

In some cases, this extended to what practitioners termed community mitosis. The term borrows from an established concept in the online community literature, where mitosis describes a mature community splitting into smaller, self-sustaining sub-groups as it grows (Iriberri & Leroy, 2009). Respondents described a distinct, harm-driven variant of the same process: rather than splitting to accommodate healthy growth, members split off into smaller, closed or private spaces to escape hostility and continue their engagement away from public view. This protected individual members, but respondents noted that it also fragmented communities, reduced the diversity of participation in the original space, and shifted the burden of safety onto smaller, less visible and often less resourced groups. Where growth-driven mitosis is generally read as a sign of a community's success, hostility-driven mitosis is better read as a sign of its failure to protect the members it already has.

Survey respondents described the same pattern independently, in similar language:

> Definitely has a chilling effect–lots of people just give up and fade away over time. More splintering into safer, smaller private groups. Contributes to collective exhaustion and burnout.
>
> - SURVEY RESPONDENT

Derailed conversations

Beyond withdrawal, respondents described misogyny actively derailing conversations in progress.

Discussions on women's achievements, safety, leadership or gender equality were consistently redirected towards the behaviour of the person raising them, or towards familiar counter-narratives, rather than the substance of what was being discussed. Even where a conversation continued after trust and relational repair by moderators, its original purpose was frequently difficult to retrieve.

Real-world harms

Some respondents also described misogyny escalating beyond the platform itself, into forms of harm that reach into members' offline lives, from safety concerns to the erosion of trust in the spaces and organisations meant to protect them.

Practitioners described this as a particularly difficult governance challenge: platform moderation tools are built to act on content, not to respond to harm that has moved beyond the platform's boundaries, leaving practitioners to manage risks their tools were never designed to address.

One social media manager at a major superannuation fund described her team's response after a direct threat: they moved from signing off responses with their first names to gender-neutral aliases, and built a close working relationship with the organisation's physical security team.

One social media manager at a major superannuation fund described her team's response after a direct threat: they moved from signing off responses with their first names to gender-neutral aliases, and built a close working relationship with the organisation's physical security team.

> We just didn't carry the same weight or respect as we would if I signed off with my own name. [Using aliases] has sort of just neutralised it—we shrunk the target, I guess.
>
> - PRIYA
> *social media manager, superannuation sector*

Mental health issues for moderators

More attention is being paid to mental health in moderation, although much of this work tends to focus on what Sarah Roberts (2019) calls 'commercial content moderators' - the people who work at large scale and volume for social media platforms themselves. Our survey also brought to the fore the impact that the work has on people who work at smaller scales with respondents describing variously growing disillusionment, desensitisation and acceptance that this is just the price of social media.

> I feel we are providing a valuable service to the community in moderating these comments or content. It does often leave me feeling disillusioned with society though... Obviously there is progress in human history, we've come a long way, but the persistence of ignorance and encouragement by algorithms or political figures who benefit from division and populist movements, makes it feel like we're going backwards.
>
> - SURVEY RESPONDENT

Rather than being triggered by larger individual events, the moderators we spoke to described the effects of the work as a result of repeated low grade exposures that compounded over time. This is in line with research on burnout, compassion fatigue and complex post-traumatic stress disorder (C-PTSD) in the workplace (Williams & Williams, 2020). One moderator described the specific toll of seeing the same hostility recur with regularity, as part of a paid role:

> When it's your 9 to 5, and you see these comments pop up with regularity, it starts to make you jaded and cynical, because it's a reminder—you see their profile pictures and think, that looks like anyone I could walk past. These are people out there who carry these beliefs, and they're just like you and I on Facebook doing whatever. The emotional toll of that...it can be quite significant.
>
> - PAIGE
> *online engagement team leader, mental health support forum*

This sense of disillusionment was widely shared.

The most commonly reported impacts on practitioners across the survey were:

63%
Emotional fatigue or burnout

44%
Reduced enjoyment of the communities they manage

42%
Reduced trust in platforms

37%
Reduced social media use outside work

35%
Desensitisation to harmful content

28%
Anxiety or stress

Respondents connected these impacts to burnout and, in some cases, to their leaving the role altogether. This additional emotional labour fell particularly heavily on moderator leads and community managers responsible for supporting other moderators, who absorbed both their own exposure to misogynistic content and the emotional fallout experienced by the teams they led.

Respondents also described adapting their own behaviour in response to the relentlessness of misogyny, including avoiding social media outside working hours, withdrawing from public participation in the spaces they manage, hiding or muting content rather than engaging with it directly, and declining to post content they anticipated would attract hostility. These self-protective strategies, while understandable, represent a further cost of the work and can themselves shape moderation outcomes: they narrow practitioners' own participation in much the same way misogyny narrows participation for the members or followers they protect, and can mean that early warning signs go unaddressed because a practitioner has disengaged to protect their own wellbeing.

The desensitisation many respondents described warrants attention. Several reported feeling numb, or that repeated exposure had stopped affecting them the way it once did; interview participants described becoming "jaded and cynical", or needing to grow "tougher" simply to keep doing the work. This numbing, alongside the avoidance behaviours described above and the hypervigilant scanning for risk several treated as routine, maps closely onto the four symptom clusters associated with post-traumatic stress disorder: intrusion, avoidance, negative alterations in mood and cognition, and hyperarousal (Phelps et al., 2022).

Research on commercial content moderators documents comparable trauma responses, drawing parallels with emergency services and care professions like social work (Spence et al., 2023). This report cannot, and does not, diagnose individuals. But the consistency between practitioner accounts and the clinical literature on complex trauma and resilience is difficult to ignore, and strengthens the case for treating repeated exposure to misogynistic content, alongside other harmful behaviours moderators govern online, as a genuine occupational health and safety concern.

Support structures for this labour remain uneven. Almost half of respondents (49%) relied on informal support from colleagues and peers. Formal workplace support, such as an employee assistance program, was also available to 47% of respondents, while 37% accessed external mental health services (such as a private psychologist). However, 2% reported having no support available to them at all. One in five reported minimal or no organisational support, and 16% worked independently, without access to organisational structures altogether.

If community managers, moderators and social media managers constitute a front line of digital defence for Australian society against socio-cultural harms, their wellbeing should be understood as a governance concern in its own right, not an incidental cost of the work.

What makes it worse

Across both the survey and interview data sets, respondents identified a consistent set of conditions that compounded the difficulty of moderating gendered harms in their digital social spaces.

Under-resourced tools and platform abdication

Respondents were unconvinced that current moderation tools and technologies are adequate to the task, with 63% of survey respondents rating available tools as somewhat or very inadequate (against just 2% who rated them very adequate). Automated systems were noted as capable of assisting with explicit slurs and clear rule violations, but struggled with the subtlety, irony and coded language that characterised much of the misogyny experienced by practitioners.

Several respondents were blunt about who they believe is culpable:

> The biggest elephant in the room is the platforms themselves, and their complete lack of motivation to improve what appears on their platforms, because it drives further engagement. If I could wave a magic wand, I would try and make those billionaires give a damn.
>
> - ROBYN
> *agency-based community manager*

Practitioner dissatisfaction extended to platform reporting and enforcement processes themselves. Respondents frequently described content remaining online despite clear platform policy breaches, reports dismissed without explanation, limited or non-existent appeals processes, and repeat offenders facing few meaningful consequences.

Meta was singled out repeatedly, with several respondents describing misogynistic content routinely judged not to breach platform policy despite clear harm to their audiences and communities. For many respondents, platform reporting systems functioned less as a safeguard than as an unreliable final step after the substantial moderation work had already been done by practitioners themselves. Many said they "wouldn't even bother" reporting, especially if they could step in themselves.

> Facebook struggles to understand what misogyny is. It doesn't recognise that misogyny is disinformation.
>
> - SURVEY RESPONDENT

Overwhelm

Both survey respondents and interview subjects pointed to sheer volume as a compounding factor:

hostile content and behaviour spread across multiple platforms and channels simultaneously, often faster than it could be reviewed, regardless of the tools or institutional support available.

Even well-resourced teams described being unable to keep pace with the volume of activity across the full range of digital spaces they were responsible for.

Contextual complexity

Respondents also identified the contextual nature of the work itself as a compounding difficulty.

As earlier sections of this report establish, distinguishing misogyny from disagreement, and coordinated harassment from isolated incidents, relies on historical knowledge, relationship-building and cumulative pattern recognition that cannot be reduced to a checklist. One example offered was a user in a help-seeking online community who likely had a mental illness (based on community context and 'digital body language'). The community manager could not simply "ban or block" the user as they might with a known bad actor, but they must devise specialised communication strategies to counsel and guide behaviour, while maintaining care for the community collective experience.

This complexity cannot be resolved through better tools alone: it is inherent to the situated judgement the work requires.

Lack of understanding and respect

All respondents and interview subjects spoke to the lack of understanding of, and respect for, what moderation work actually involves as a compounding factor for the harms they experience and govern on behalf of others.

This extended across organisations and into public discourse. When moderation succeeds, its effects are often invisible: conflicts de-escalate before they become visible, harmful content is removed before members encounter it, and the space continues functioning without disruption.

This invisibility works against practitioners twofold: it leaves the labour underappreciated and under-respected, and it leaves policymakers and researchers focused on platform-level moderation decisions while overlooking the substantial governance work occurring behind the scenes.

The online/offline divide

Many respondents connected their difficulty in having harms recognised to an institutional 'othering' of the online spaces they govern.

This report treats that framing as an artificial divide. It positions the space practitioners govern as an abstraction, "cyberspace" or "the internet", detached from the offline lives of the people occupying it, rather than a live extension of them.

Practitioners experienced this framing as delegitimising. Several described institutions, including law enforcement, treating threats and abuse encountered through their work as somehow less serious or less actionable because they occurred through a screen rather than face to face, despite carrying comparable psychological and, at times, physical risk. This echoes the pattern of institutional betrayal discussed elsewhere in this report: a failure to take the harm seriously in the first place, compounded by the further harm of not being believed once it is raised.

The same framing shapes regulation. Respondents observed that policy responses to online harms often relabel existing protections with the word "digital," rather than genuinely accounting for how these harms work, producing duplicative frameworks that add complexity without adding protection. Practitioners (as 'micro-regulators') are left navigating a patchwork of policy built on the assumption that online and offline harm differ in kind, rather than in medium alone.

This abstraction cuts both ways. It also makes it easier to dehumanise the people participating in, and governing, digital spaces, reducing someone to a username or a profile picture rather than a person with a name and a life beyond the screen. Respondents connected this directly to the behaviour they were required to moderate, noting that abusive users often seemed unable to conceive of the person on the other end of a comment as a person at all.

Several described a similar failure of imagination in how their own labour was perceived by employers and the public.

Participants also spoke to AI and the rise of synthetic culture is compounding this problem, from the form abuse itself now takes to a growing assumption that the humans moderating a space are themselves ‘bots’, and therefore not owed the basic recognition extended to a person.

Several interviewees suggested a public awareness campaign, for example, a ‘day in the life of a moderator’, to show the work involved, the abuse they encounter, and the human behind the screen. Making the labour visible in this way could build public understanding directly, and give institutions, including law enforcement and employers, a concrete reference point for a form of work that is otherwise easy to imagine as automated, anonymous or already handled by someone else.

Structural disempowerment

The organisational separation some participants described is better understood as structural disempowerment than as a simple staffing or workflow issue.

This follows a pattern with a longer history. Cultural-historical research (Paech, 2025) traces how online spaces moved from independently built communities towards corporate platforms. Technical agency was progressively decoupled from socio-cultural practice. Platform architectures absorbed control over infrastructure. Community practitioners were repositioned as culturalists, responsible for shaping social and affective conditions, but increasingly excluded from the technical and financial decisions that shape those conditions. This report's findings suggest that the same decoupling is still unfolding, now within individual organisations rather than only across the platform economy.

One respondent's own team structure illustrates the pattern directly:

> We have a strong community management focus, so we don't sit within the marketing or the brand team, and we don't look after publishing or creating content. Our role is purely managing interactions and answering questions.
>
> - PRIYA
> *social media manager, superannuation sector*

Participants in this research descibed this problem occuring repeatedly in workplaces. The people who write and post content — often marketing or communications staff — work separately from the people who deal with the fallout, such as community managers, customer service teams or moderators, with minimal contact or influence between them. Moderators who understood exactly which content triggered hostility, and could anticipate it, had no ability to change what was posted in the first place. This structural marginalisation compounds every other harm identified in this report: practitioners are left managing the downstream consequences of decisions made upstream, by people who may never see, or be asked to account for, the hostility those decisions generate.

For practitioners working outside formal organisations, structural disempowerment took a different form: exclusion from training, resources and peer support. Respondents in this position described building the knowledge and confidence to moderate effectively largely alone, without access to the professional development, mentoring or community of practice available to their organisationally embedded counterparts. Some actively resisted this isolation, forging informal peer networks and knowledge exchange regardless. But respondents were clear that these networks formed despite the conditions of their work, not because of them: the systems and structures around them were not designed to support collective learning or mutual support, and in practice tended to further isolate practitioners from one another.

Taken together, these findings paint a problematic picture: practitioners' capacity to implement effective practices is stymied by structual design or constraints.

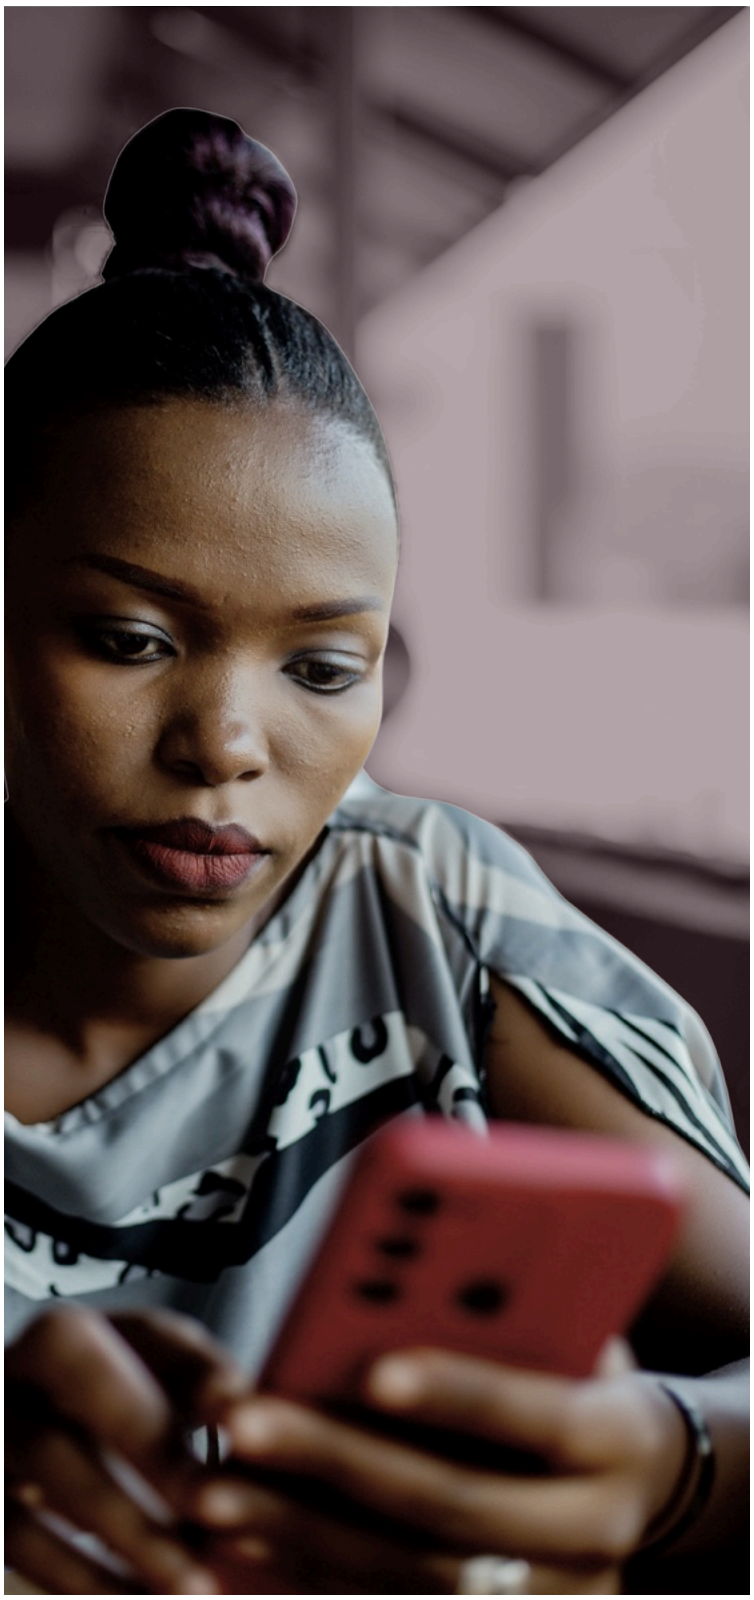

Institutional betrayal

The institutional gaps described earlier in this report were compounded, for some respondents, by the gendered nature of the labour itself.

Self-identified respondents were more likely to describe the personal and emotional toll of the work, yet institutional support structures, where they exist, rarely account for this dimension explicitly, tending instead towards generic wellbeing provisions that do not reflect the specific nature of gendered hostility.

Read together, these findings describe a pattern of institutional betrayal rather than simple neglect. Institutional betrayal, as defined by Freyd and colleagues, occurs when an institution fails to prevent harm to people who depend on it, or fails to respond supportively once that harm is disclosed, compounding the original harm with a second, institutional one (Smith & Freyd, 2014; Freyd & Birrell, 2013). Practitioners in this research described both forms: harm they were left to absorb because their organisation had not prepared for it, and harm compounded by organisations that minimised, delayed or reframed it once it was raised.

Under-resourcing can be solved with a budget line. Institutional betrayal requires something closer to what Freyd (2014) calls 'institutional courage': transparency, accountability and a willingness to prioritise the people an institution depends on over its own reputation. Some respondents noted that their organisations were capable of this, but that action depended heavily on individual leadership whims.

One consultant with an extensive background in the non-profit and health charity sectors described a case from her own career that illustrates the pattern in full. After banning a user who had subjected a vulnerable community member to sustained abuse, following what she described as hundreds of prior reported incidents, she faced months of escalating complaints from the banned user directly to senior leadership. Her organisation reopened a closed case, appointed a male manager with no moderation experience to investigate her own decision, and disclosed to her, alone, that the individual had been attempting to identify her contact details, without escalating the safety risk internally. She resigned soon after.

> “...even in organisations that are supposedly about duty of care, [they] don’t give a shit. They’re not prepared to back you, they’re not prepared to put things in place, and you end up just having to leave, because that’s the only way to escape it.
>
> - CLAIRE
> *community management consultant, non-profit and health sectors*

Other respondents described the difference clear institutional backing makes. One practitioner, whose organisation's managing director had explicitly endorsed his team's authority to block, ban and remove without needing to justify each decision, described feeling “very empowered to be able to make those moderation decisions without having to be questioned about it, or fear of repercussions,” and noted this did not appear to be the norm across the sector (Tom, corporate affairs practitioner, retail sector). The contrast between these two accounts is, in itself, a finding: institutional courage is possible, but it remains a matter of individual leadership rather than structural guarantee.

A systemic problem

> The findings in this report describe practitioners governing symptoms of a problem whose causes sit beyond the audiences and communities they govern, or any single platform.

As established earlier in this report, respondents consistently linked the volume and intensity of misogyny they encountered to the broader social and political climate:

a climate in which prominent media figures and political leaders routinely deploy misogynistic rhetoric, in which algorithmic amplification rewards division, and in which online spaces increasingly serve as a proving ground for narratives that then migrate into offline life, and back again. Online misogyny is not a phenomenon contained by, or created by, the platforms on which it appears. It is a reflection of attitudes and behaviours that also exist offline, expressed through infrastructure that often concentrates and accelerates it.

This matters because online harms still aren't treated with the seriousness of their offline equivalents. According to participants, most of their time is spent managing behaviour that sits below the threshold of existing legal or platform policy frameworks: e.g. subtle sexism, dog whistles, exclusionary conduct, coordinated harassment, and cumulative patterns of hostility. Regulatory attention to online harms in Australia and elsewhere has, understandably, focused heavily on protecting children; the governance of gendered hostility affecting adult women receives comparatively little attention, despite its scale and its documented effects on participation, wellbeing and safety. Survey data reflects this gap in practice. Just 41% of respondents said they were aware of relevant legislation, such as Australia's Online Safety Act, and applied it in their decisions. A further 34% said they were aware of the legislation but that it didn't feel relevant day-to-day, 23% said they weren't really aware of it, and 23% didn't believe it applied to their context (respondents could select more than one option). This points to a legislative framework that still has much work to do.

For many respondents, this gap was compounded by a further, more diffuse form of institutional betrayal: watching successive governments fail to develop regulation with 'teeth', or force platforms to do the right thing.

Respondents described exhaustion with a longstanding "performative" pattern: platforms are repeatedly pushed, through regulation, public pressure or direct advocacy, to take gendered harm seriously, and repeatedly don't, while the government bodies meant to hold them accountable produce little beyond further consultation. Practitioners depend on both platforms and governments to make their community governance work possible. When neither delivers, institutional betrayal isn't confined to a single employer; it operates at the level of the whole system practitioners are asked to trust.

Regulation, where it exists, establishes the outer boundary of acceptable conduct. Practitioners are responsible for governing the much larger grey area beneath that boundary, largely without formal legal or policy support, using judgement, relationship-building and community knowledge developed over years of practice. Addressing misogyny in these online spaces therefore requires action at both levels: continued investment in the frontline governance practices this report documents, and broader social, political and regulatory change that addresses the conditions producing the behaviour in the first place.

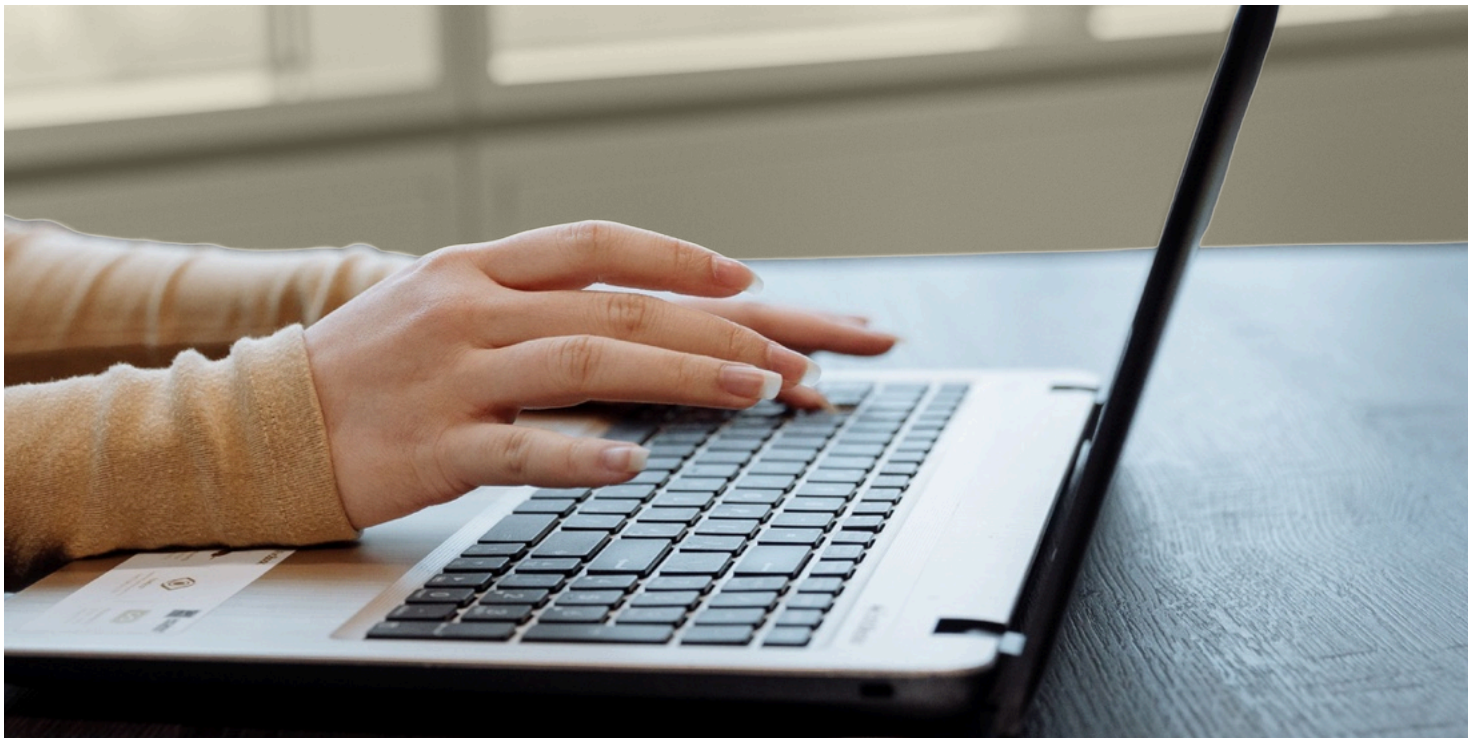

Conclusion

This report makes visible a form of labour that has, until now, remained largely unrecognised: the everyday work of community managers, moderators and social media managers in identifying, interpreting and responding to misogyny within the digital social spaces they govern.

The findings show that misogyny in these online social spaces is routine rather than exceptional, cumulative rather than isolated, and inseparable from the broader social and political conditions in which it occurs. They show that practitioners do not simply enforce platform rules; they interpret behaviour within its social context, distinguish disagreement from hostility, weigh institutional and personal risk, and choose between enforcement, education and restorative practice according to circumstances that platform systems are neither designed nor resourced to assess. And they show that this work carries a significant personal cost, borne disproportionately by the practitioners performing it, with support structures that remain informal, uneven and often absent altogether.

Read together, the findings answer the questions this report set out to address. Online social spaces become safer for women and diverse voices under identifiable conditions:

- active and consistent community stewardship is present
- community stewards implement holistic moderation that includes pre-emptive and responsive measures, and a combination of educative, restorative and punitive interventions
- community stewards have agency (and skills) to act on contextual judgement rather than content rules alone
- institutions back visible, proactive governance instead of defaulting to a lighter touch
- safety work extends beyond removing harmful content to supporting the people it targets

Some of the techniques that produce this safety already exist in practice: contextual forensics that reads behaviour in context rather than isolation, graduated and restorative responses that address harm without excluding good-faith participants, bystander activation that distributes the labour of pushback, and record-keeping that turns individual vigilance into institutional memory. None require new technology — they require practitioners who are trusted, resourced and supported to use the judgement they already have.

Three findings in particular warrant the attention of platforms, policymakers, funders and the organisations that rely on this work:

The psychosocial risks of moderation work require formal recognition as an occupational health and safety issue, with obligations that sit with employers and platforms, not only with the individual practitioners who absorb the impact.

Meaningful accountability for misogynistic behaviour requires coordinated action across government, business and civil society. Informal, practitioner-led governance cannot indefinitely substitute or gap-fill for structural accountability. Improvement demands institutional courage to prioritise people over reputational or commercial interest.

An embedded workforce, largely untapped organisations, governments and safety actors have yet to fully recognise or invest in this asset. They could better leverage these practitioners and the insights they hold, rather than defaulting to generic guidance or technology solutionism.

Community managers, moderators and social media managers are already doing sophisticated, effective governance work, often without formal recognition, training or support commensurate with its difficulty. The task now is not to invent new solutions, but to recognise, resource and scale the ones already in practice.

Next steps

This report shows that effective responses to misogyny in digital social spaces exist, embedded in the everyday practice of community managers, moderators and social media managers. The task now is to better recognise, resource and scale that expertise, and to address the social and political conditions that sit beyond any individual practitioner's control.

Asked directly what role practitioners should play in any future digital duty of care framework, one agency-based community manager was unequivocal:

> We will be the people that are probably implementing this in reality. We are the most embedded on that front line. You're never going to get buy-in from platforms, so it will be incumbent on organisations to buy in, and us to make it happen. Maybe they should get community managers to draft it.
>
> - ROBYN
> *agency-based community manager*

Community managers, moderators and social media managers are uniquely positioned to shape this next phase of work. They hold contextual knowledge, relationship-based authority and lived operational experience. Platform trust and safety teams cannot readily replicate this, shaped as they are by commercial and legal priorities. Nor can government bodies, which operate at a policy remove from the day-to-day realities of this work. Solutions designed and led by practitioners embedded on front lines, grounded in practice that is already working, are more likely to be adopted, trusted and sustained than approaches imposed from outside the audiences and communities they are meant to protect.

Building on the findings of this report, Digital Rights Watch and Australian Community Managers propose the following areas for further investment and research.

A national-scale survey	This study's findings, drawn from a qualitative and preliminary quantitative sample, warrant testing across a nationally representative sample of practitioners, to establish the prevalence, cost and distribution of these findings with the statistical weight required to inform policy and funding decisions.
Formal training and standards	National standards and training in community management and moderation practice, for practitioners, and for organisations that engage individuals or services offering this labour (or perform it themselves).
Occupational health and safety guidance	Development of specific OH&S wording and guidance for moderation as an occupation, recognising its psychosocial risks in the same way other high-exposure occupations are recognised.

A consultant with extensive experience in the non-profit sector drew direct comparison to occupations where this recognition already exists:

> I'd love to see more case studies about how you protect the people who are user-facing around harassment and abuse, kind of like the policies we've seen for the bus driver or shop keeper, with zero tolerance for abuse. We need that sort of protection extended into online labour.

- CLAIRE
community management consultant, non-profit and health sectors

A legislative and regulatory audit	A structured review of existing legislation relevant to online gendered harms, to identify genuine gaps rather than duplicating protections that already exist under a "digital" label, and to clarify where responsibility sits across platforms, employers and government.
A public awareness campaign	A public-facing campaign that makes moderation work visible, reframing moderators and community managers as people carrying out a form of frontline digital labour, and building public understanding of the psychosocial risks involved.
Research into engaging men as active bystanders	Further research into how men, both as community members and as practitioners, can be supported and equipped to recognise and call out misogyny when they see it, building on this report's preliminary finding that professional experience, not gender alone, predicts misogyny literacy. As one survey respondent put it: "no amount of law or tools will change it until men start challenging other men."
Digital civics education	Investment in school-based education about being a good digital citizen and/or community member online, the role of moderators and the human costs of harms, addressing the conditions producing this behaviour before people reach adulthood.

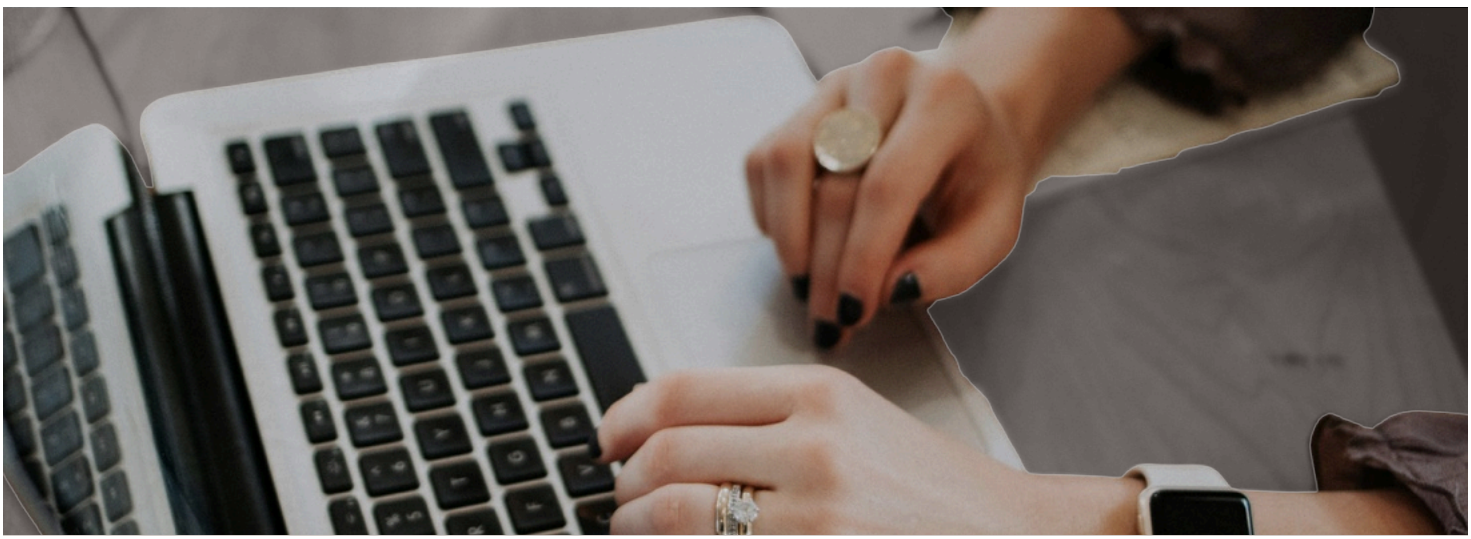

References

Anderson, B. (1983). *Imagined communities: Reflections on the origin and spread of nationalism*. Verso.

Bingham, A. J. (2023). From data management to actionable findings: A five-phase process of qualitative data analysis. *International Journal of Qualitative Methods*, 22, 1–11. https://doi.org/10.1177/16094069231183620

Blanchard, A. L. (2008). Testing a model of sense of virtual community. *Computers in Human Behavior*, 24(5), 2107–2123.

Blanchard, A. L. (2024). Revising sense of community to understand typical and extremist virtual communities. *Group Processes & Intergroup Relations*.

Blanchard, A. L., & Markus, M. L. (2004). The experienced "sense" of a virtual community: Characteristics and processes. *The DATA BASE for Advances in Information Systems*, 35(1), 65–79.

Blanchard, A. L., Welbourne, J. L., & Boughton, M. D. (2011). A model of online trust: The mediating role of norms and sense of virtual community. *Information, Communication & Society*, 14(1), 76–106.

Freyd, J. J. (2008). *The psychology of betrayal trauma: Memory, health, and gender* [Conference presentation]. Thompson Hall Science and Mathematics Seminar.

Freyd, J. J. (2014, November 9). Use science as tool on campus sexual assault [Op-ed]. *The Register-Guard*, H4.

Freyd, J. J., & Birrell, P. (2013). *Blind to betrayal: Why we fool ourselves we aren't being fooled*. Wiley.

Granovetter, M. S. (1973). The strength of weak ties. *American Journal of Sociology*, 78(6), 1360–1380.

Grüning, D., & Kamin, J. (2025). Prosocial design in trust and safety. *arXiv*. https://arxiv.org/abs/2506.12792

Hartsock, N. C. M. (1983). The feminist standpoint: Developing the ground for a specifically feminist historical materialism. In S. Harding & M. B. Hintikka (Eds.), *Discovering reality: Feminist perspectives on epistemology, metaphysics, methodology, and philosophy of science* (pp. 283–310). D. Reidel.

Iriberri, A., & Leroy, G. (2009). A life-cycle perspective on online community success. *ACM Computing Surveys*, 41(2), Article 11.

Jones, Q. (1997). Virtual-communities, virtual settlements and cyber-archaeology: A theoretical outline. *Journal of Computer-Mediated Communication*, 3(3).

Koh, J., Kim, Y. G., & Kim, Y. G. (2003). Sense of Virtual Community: A Conceptual Framework and Empirical Validation. International Journal of Electronic Commerce, 8(2), 75–94.

McMillan, D. W., & Chavis, D. M. (1986). Sense of community: A definition and theory. *Journal of Community Psychology*, 14(1), 6–23.

Newton, C. (2019, February 25). The trauma floor: The secret lives of Facebook moderators in America. *The Verge*.

Paech, V. (2025). The hidden architects of the internet: Rebecca Newton and the making of online community management. *Internet Histories*, 9(4), 364–382.

References

Paech, V. (2025). 'Online Community Managers: Learning from the Original Trust and Safety Practitioners'. In *Trust, safety, and the internet we share: Multistakeholder insights*. M. L. Daniel, A. Menking, M. T. Savio, & J. Claffey (Eds.), Taylor & Francis.

Perrigo, B. (2022, February 14). Inside Facebook's African sweatshop. *TIME*.

Phelps, A. J., et al. (2022). Australian guidelines for the prevention and treatment of posttraumatic stress disorder: Updates in the third edition. *Australian & New Zealand Journal of Psychiatry*, 56(3), 230–247.

Pinchevski, A. (2023). Social media's canaries: Content moderators between digital labor and mediated trauma. *Media, Culture & Society*, 45(1), 212–221.

Roberts, S. T. (2019). *Behind the screen: Content moderation in the shadows of social media*. Yale University Press.

Smith, C. P., & Freyd, J. J. (2014). Institutional betrayal. *American Psychologist*, 69(6), 575–587.

Spence, R., Bifulco, A., Bradbury, P., Martellozzo, E., & DeMarco, J. (2023). The psychological impacts of content moderation on content moderators.

Sugiura, L. (2021). The Incel Rebellion: The Rise of the Manosphere and the Virtual War Against Women. (Emerald Studies in Digital Crime, Technology and Social Harms). Emerald Publishing Limited.

Williams, S. D., & Williams, J. (2020). Posttraumatic stress in organizations: Types, antecedents, and consequences. *Business & Society Review*, 125, 23–40.

Wispé, L. G. (1972). Positive forms of social behavior: An overview. *Journal of Social Issues*, 28(3), 1–19.

About the authors

Dr. Venessa Paech

Dr. Venessa Paech is a digital sociologist, practitioner-researcher and Co-Founder and Director of Australian Community Managers (ACM), the national centre of excellence for online community management. Established in 2009, ACM is now home to a 6,000-member practitioner community and provides research, training, consultancy and liaison with government and civil society. She has led community and moderation strategy for organisations including Lonely Planet, REA Group, Envato and Australia Post, and consults on community design, governance and moderation for clients spanning media, government, education and the non-profit sector. In 2023 she founded All Things in Moderation (ATIM), a global conference on the practice of moderation, and from 2011 to 2022 she organised and hosted Swarm, Australia's national conference for online community managers. Venessa teaches online community management at the University of Sydney and the University of Melbourne, serves as an expert legal witness on social media matters, and co-authored Australia's national Code of Ethics for community and social media professionals.

Dr. Jennifer Beckett

Dr. Jennifer Beckett researches across three core areas: Online hostility, platform governance, trust and Safety, and moderation. This work intersects around trauma and resilience of digital workers. She has professional experience as an online and social media producer at the ABC and as a volunteer moderator for various interest-based communities. Jennifer is a founding member of Australian Community Managers and a co-convenor of All Things in Moderation. She developed and teaches a subject on Online Community Management at the University of Melbourne and works extensively with ACM in the development of guidelines and training. She is currently a CI on the ARC Discovery Project "Addressing Online Hostility in Australian Digital Cultures" (DP230100870)".

Appendix A

Survey questions

- Which of the following best describes your primary role? (select one) (Professional community manager; Professional social media manager; Professional moderator; Volunteer community manager; Volunteer social media manager; Volunteer moderator/admin; Other)
- How long have you been doing this work? (Less than 1 year; 1–3 years; 4–7 years; 8+ years)
- What type of space(s) do you primarily manage? (select all that apply) (News or media organisation; Entertainment or fan community; Corporate/brand community; Advocacy, activist, or feminist community; Gaming; Sport; Local government or community services; Other)
- Where are you primarily based? (Australia; New Zealand; Southeast Asia; Other APAC; Outside APAC)
- What platform(s) do you work on most often? (select all that apply) (Facebook; Instagram; WhatsApp; TikTok; Reddit; Discord; YouTube; Dedicated community platform, e.g. Circle, Discourse, Mighty Networks; Other)
- How frequently do you encounter misogynistic content or behaviour in the spaces you manage? (Rarely; Occasionally; Regularly; It's a constant presence)
- How does misogyny typically show up in your space(s)? (select all that apply) (Targeted abuse directed at specific individuals; Generalised sexist commentary or "jokes"; Coordinated pile-ons or harassment campaigns; Visual content, e.g. memes, AI images; Subtle or coded content that's hard to classify; Content directed at me personally as a moderator/admin/manager; Other)
- How confident do you feel in identifying misogynistic content? (Very confident; Somewhat confident; Not very confident; Unsure)
- Talk a bit about how you detect it. Are there patterns or signs you look for? What makes it easier or harder to spot? (open text)
- When you encounter misogynistic content, how do you usually decide how to respond? (select all that apply) (Follow formal organisational policies; Apply my own judgement; Consult a colleague or supervisor; Refer to community guidelines; Use a combination of policy and personal judgement; Other)
- What strategies or practices do you use to manage this type of behaviour, especially those things you've figured out through experience? (open text)
- How adequate are the tech tools available to you (platform features, reporting systems, automation) for managing misogynistic content? (Very adequate; Adequate; Somewhat inadequate; Very inadequate)
- What's the biggest gap or issue with your tools or processes? (open text)
- Are you aware of relevant legislation or regulation, such as Australia's Online Safety Act, that applies to your work? (select all that apply) (Yes, and I apply it in my decisions; I'm aware but it doesn't feel relevant day-to-day; I'm not really aware of it; I don't believe it's applicable to my context)
- How has moderating misogynistic content affected you personally? (select all that apply) (Emotional fatigue or burnout; Desensitisation to harmful content; Anxiety or stress; Reduced enjoyment from your community; Using social media less; Reduced trust in platforms; Hasn't really impacted me that I'm aware of; Other)
- How much support do you receive from your organisation when managing misogynistic content? (Strong support; Some support; Minimal support; No support; I work independently, no organisation)
- How has this misogynistic content or behaviour affected your community/users/audience/members? (open text)
- Do you have access to wellbeing support or mental health resources related to this work? (select all that apply) (Yes, formal support is available through my job; Yes, I can access formal support via non-work related external services, e.g. a psychologist, Lifeline; Yes, I have informal support from colleagues or peers; I have no support accessible to me; I haven't needed it)
- Are there specific content or activities in your space that you've found more likely to attract misogynistic commentary and behaviour? (open text)
- What single change would make the biggest difference to how misogyny is managed in your space? (open text)
- Would you be willing to participate in a follow-up interview (approx. 60–90 minutes, conducted via Zoom)? (Yes; No)
- Share your email if you answered YES and are open to a follow-up interview. (open text)

Appendix B

Interview questions (semi-structured)

Background:

- Can you describe your audience/account/channel/community and your role in looking after it?
- How long have you been doing this work, and what drew you to it originally?
- Has your relationship with the role changed over time?

Interpretation, decision-making and tacit knowledge:

- How do you draw lines around behaviour in your online space?
- How do you know someone has crossed them?
- When you encounter something that feels harmful or borderline, what's your process for deciding how to respond?
- How much of that is guided by formal policy versus your own judgement?
- What informal strategies or practices have you developed over time that you wouldn't necessarily find written down anywhere — things you've figured out through experience?
- What tools do you use to help you here, and how effective are they?

Encountering and identifying misogyny

- How much misogynistic content would you say you come across in the spaces you look after?
- Has this increased or decreased recently, and what do you think is behind that?
- How does misogyny typically show up? Are there patterns (actors involved, timing, tactics, targets) that you've come to recognise?
- Where does it get complicated? Are there behaviours that are hard to classify, or where you're uncertain whether something crosses the line?
- How does it affect your community or followers?
- Do you have any specific strategies for dealing with misogynistic content and its fallout? Could you outline them for me?

For content-led communities or social media spaces

- What kinds of content that you post are most likely to attract misogynistic commentary and behaviours?
- When you know this content is likely to attract misogynistic responses, how do you prepare? What do you do differently?
- Are there any different policies or procedures in place in your workplace?

Community dynamics

- How do you navigate community members who push back on your moderation decisions, for example when misogynistic behaviour is framed as “just a joke” or free speech?
- Are there ways your community itself has developed protective norms or taken on a role in calling out harmful behaviour? How do you foster or support that?

Appendix B

Institutional and structural conditions

- What bigger-picture factors (resourcing, leadership support, platform tools, legal constraints) make this work easier or harder for you?
- How does the law shape your day-to-day decisions, if at all? Are you aware of regulations like the Online Safety Act? Do they feel relevant to the work you're actually doing?

Emotional and relational labour

- What does it feel like to moderate misogyny, particularly if it's directed at you, community members you know, or when it's sustained and repeated? How does it impact you (for example, your wellbeing or relationship to the platforms you work on)?

Improving things

- What do you wish others (for example platforms, employers or policymakers) better understood about what this work actually involves?
- What responsibilities do you think people have in fighting misogyny online? (Meaning everyday people, not just policymakers.)
- If you could create a curriculum for people, what would you put in it, and who would you get to deliver it?
- What would make you feel safer doing this work?
- What would you change about how misogyny in your community, and more broadly, is moderated (for example the tools available to you, the structures you work within, or the broader environment)?
- Which of these changes do you think is the most important? How would you prioritise them?
- Have any tools been introduced that have actually made your work harder?
- The government is considering introducing a Digital Duty of Care legislation. How do you interpret the idea of a Digital Duty of Care? Who has it, and what should be included in that legislation?

Closing

- Is there anything else you would like to raise or mention?

Appendix C

Interview participant summary (anonymised)

All interview participants were offered a pseudonym; all six elected to use one. Organisation names and other identifying operational detail have been removed or generalised. Years of experience refer to time in community management, moderation or social media management roles broadly, not necessarily in the participant's current position.

- **Paige** — team leader responsible for online engagement and moderation at a national mental health peer-support organisation, covering both peer-support forums and social media channels. Approximately five years' experience in community moderation, clinical background.
- **Frankie** — social media and community management practitioner with more than two decades' experience across public institutions, higher education, financial technology and, most recently, superannuation.
- **Tom** — community and corporate affairs practitioner responsible for social media issues management and moderation governance at a large Australian retail organisation. Approximately 10 years in the role.
- **Priya** — social media manager leading a small team at a major Australian superannuation fund, responsible for community interactions and moderation across owned social channels.
- **Robyn** — agency-based community manager supporting multiple clients across sectors, based regionally. Approximately five years' experience, having moved into community management from a prior career.
- **Claire** — community management consultant with extensive experience in the non-profit and health charity sectors, specialising in peer-support communities and “communities of circumstance.”

COPYRIGHT 2026
AUSTRALIAN COMMUNITY MANAGERS AND DIGITAL RIGHTS WATCH